

Package ‘GHAgricProductivityLab’

November 29, 2025

Type Package

Title Agricultural Productivity in Ghana

Version 0.0.0.9000

Author Francis Tsiboe [aut, cre] (<<https://orcid.org/0000-0001-5984-1072>>)

Maintainer Francis Tsiboe <ftsiboe@hotmail.com>

Contributor -

Reviewer -

Creator Francis Tsiboe

Description Provides tools and datasets for investigating the drivers of agricultural production shortfalls in Ghana.

It compiles research studies that examine farmer-specific and institutional factors to assess whether inefficiencies arise from technical inefficiency, technology gaps, or both. The package offers empirical evidence to inform policy discussions and interventions aimed at improving agricultural productivity, particularly in contexts with limited access to modern technologies.

License GPL-3 + file LICENSE

URL <https://github.com/ftsiboe/GHAgricProductivityLab>

BugReports <https://github.com/ftsiboe/GHAgricProductivityLab/issues>

Encoding UTF-8

Roxygen list(markdown = TRUE)

RoxygenNote 7.3.3

VignetteBuilder knitr

Depends R (>= 4.1.0)

Imports MatchIt, data.table, haven, stats, sfaR, piggyback

Remotes github::dylan-turner25/rfcip

Suggests knitr, crayon, tidyverse, withr, rmarkdown, lavaan, quadprog, emdbook, cobalt, cli, dplyr, doBy, stringr, MASS, Matrix, corpcor, testthat (>= 3.0.0), car, frontier, micEcon, rgenoud, purrr, randomForest, CBPS, dbarts, optmatch, Matching

LazyData true

Cite-us If you find it useful, please consider staring the repository and citing the following studies

- Tsiboe F. (2021). Chronic sources of low cocoa production in Ghana, new insights from meta-analysis of old survey data.
Agricultural and Resource Economics Review.50(2) 226-251.
- Tsiboe, F., Egyir, I. S., & Anaman, G. (2021). Effect of fertilizer subsidy on household level cereal production in Ghana.
Scientific African, 13, e00916.
- Tsiboe F. (2022). Nationally Representative Farm/Household Level Dataset on Crop Production in Ghana from 1987-2017.

R topics documented:

compute_jenks_binned_shares	2
covariate_balance	3
draw_matched_samples	4
draw_msf_estimations	6
draw_msf_summary	9
get_crop_area_list	11
get_household_data	11
harmonized_data_prep	12
match_sample_specifications	13
msf_workhorse	14
run_only_for	17
sf_workhorse	18
study_setup	21
treatment_effect_calculation	22
treatment_effect_summary	24
write_match_formulas	25

Index

26

compute_jenks_binned_shares
Compute Jenks-Binned Shares by Aggregation Groups

Description

This function bins a numeric variable using Jenks (Fisher) natural breaks and computes the relative share of *counts* and/or *weights* within each aggregation group across the binned ranges.

It is useful for summarizing the distribution of a variable (e.g., dealer density, yield, acreage, distances) across spatial or categorical groups.

Usage

```
compute_jenks_binned_shares (
  dt,
  output_variable,
  aggregation_points,
  jenks = NULL,
  jenks_number = NULL,
  weight_variable = NULL
)
```

Arguments

<code>dt</code>	A <code>data.table</code> containing the data.
<code>output_variable</code>	Character name of the numeric column to bin.
<code>aggregation_points</code>	Character vector of grouping variables (e.g., "region", "district", "grid_id").
<code>jenks</code>	Optional numeric vector of breakpoints. If <code>NULL</code> , Jenks breaks are computed automatically.
<code>jenks_number</code>	Integer number of Jenks classes to compute when <code>jenks</code> is <code>NULL</code> .
<code>weight_variable</code>	Optional character name of a weighting column. If <code>NULL</code> , a weight of 1 is assumed for each row.

Details

Internally the function:

1. Computes Jenks natural breaks (if not provided)
2. Creates `count` and `weight` totals within each bin by group
3. Merges with total counts/weights per group
4. Computes shares

Value

A `data.table` with:

- grouping variables
- `binned_range_name` - character label of the Jenks bin
- `binned_range_level` - numeric factor level of the bin
- `estimate_count` - share of counts within the group
- `estimate_weight` - share of weights within the group

See Also

Other helpers: `get_crop_area_list()`, `harmonized_data_prep()`

`covariate_balance` *Compute Covariate Balance Summaries Across Matching Specs*

Description

Compute Covariate Balance Summaries Across Matching Specs

Usage

```
covariate_balance(match_specifications, matching_output_directory)
```

Arguments

`match_specifications`
 data.frame/data.table with at least columns: `boot`, `ARRAY`, and the spec fields stored in RDS (e.g., `method`, `distance`, `link`).
`matching_output_directory`
 Directory containing RDS files named as "0001.rds", "0002.rds", etc.

Details

Reads each matching result RDS (expected to contain `m.out` and `match_specifications`), extracts balance via `cobalt::bal.tab`, reshapes, computes a composite balance $rate = \text{mean}((Diff - 0)^2, (KS - 0)^2, (V_{Ratio} - 1)^2)$, and averages by `ARRAY`, `method`, `distance`, `link`, `sample`.

Value

A list with:

<code>rate</code>	data.frame of mean balance metrics by spec (Adj sample only) with a composite rate.
<code>bal_tab</code>	long-format balance table per covariate/stat/sample/spec.

See Also

Other matching: `draw_matched_samples()`, `match_sample_specifications()`, `treatment_effect_treatment_effect_summary()`, `write_match_formulas()`

`draw_matched_samples`

Draw matched samples (stratified bootstrap + matching)

Description

Performs stratified resampling at the Surveyx, EaId level (excluding EaIds listed for the current bootstrap ID) and computes adjusted sampling weights used for matching. Then fits a matching model per `match_specifications[i,]` using **MatchIt**, with exact matching on `Emch` and `distance/link` from `match_specifications`.

Usage

```
draw_matched_samples(
  i,
  data,
  match_variables_exact,
  match_variables_scaler,
  match_variables_factor,
  match_specifications,
  sample_draw_list,
  verbose = FALSE
)
```

Arguments

i	Integer index selecting the row of match_specifications to use.
data	A data.frame/data.table containing at least: Surveyx, EaId, HhId, Mid, unique_identifier, Weight, Treat, plus columns named in Emch, Scle, Fixd.
match_variables_exact	A character vector of variable names to be included in the exact match component (e.g., "gender", "region").
match_variables_scaler	A character vector of continuous or scalar variable names to include as covariates in the general match formula (e.g., "age", "income").
match_variables_factor	A character vector of factor variable names to be included as dummy-coded categorical covariates in the general match formula.
match_specifications	Data frame of matching specifications with columns boot, method, distance, and optionally link.
sample_draw_list	Data frame where column ID identifies the bootstrap draw, and each remaining column corresponds to a Survey containing the sampled EaId.
verbose	Logical; if TRUE, prints the chosen matching method and timing. Default FALSE.

Details

Adjusted weights are computed as $pWeight = Weight \times (alloc/allocj)$, where alloc is the pre-exclusion sum of Weight by Surveyx, EaId and allocj is the post-exclusion sum.

Exact matching is performed on variables in Emch. The distance model formula is constructed as Treat ~ Scle + Fixd. When distance == "glm", m.order = "largest" is used; otherwise "closest".

Value

A list with:

```
match_specifications The selected row from match_specifications.
m.out The matchit object.
md Matched data: Surveyx, EaId, HhId, Mid, unique_identifier, weights, pWeight.
df The analysis data with adjusted weights: Surveyx, EaId, HhId, Mid, unique_identifier,
pWeight.
```

See Also

Other matching: [covariate_balance\(\)](#), [match_sample_specifications\(\)](#), [treatment_effect_ca](#)
[treatment_effect_summary\(\)](#), [write_match_formulas\(\)](#)

draw_msf_estimations

Run a single MSF draw and summarize stochastic frontier results

Description

Performs one iteration of multi-stage stochastic frontier (MSF) estimation for a given draw. The function (i) filters the analysis sample using a draw-specific exclusion list, (ii) calls `msf_workhorse()` to estimate production, inefficiency, and risk components by survey, and (iii) optionally builds disaggregated efficiency score summaries.

Usage

```
draw_msf_estimations(
  draw,
  drawlist,
  surveyy = FALSE,
  data,
  output_variable,
  input_variables,
  inefficiency_covariates = NULL,
  risk_covariates = NULL,
  weight_variable = NULL,
  production_slope_shifters = NULL,
  intercept_shifters = NULL,
  f,
  d,
  identifiers,
  include_trend = FALSE,
  technology_variable = NULL,
  matching_type = NULL,
  adoption_covariates = NULL,
  intercept_shifters_meta = NULL,
  disagscors_list = NULL,
  match_specifications,
  match_specification_optimal,
  match_path
)
```

Arguments

<code>draw</code>	Integer (or numeric coercible to integer). Draw iteration index. Typically 0 is used for the full baseline sample and positive integers for resampled or matched draws.
<code>drawlist</code>	A <code>data.frame</code> or similar object containing exclusion information by draw. Rows correspond to draws, with column <code>ID</code> identifying the draw and subsequent columns containing <code>EaId</code> values to remove from <code>data</code> for that draw.
<code>surveyy</code>	Logical; if TRUE, uses the survey labels stored in <code>data\$Surveyx</code> . If FALSE (default), all observations are assigned to a synthetic survey labeled "GLSS0" for estimation.

<code>data</code>	A <code>data.frame</code> or <code>data.table</code> containing the estimation sample. Must include, at minimum, the columns referenced in other arguments (e.g., <code>output_variable</code> , <code>input_variables</code> , <code>identifiers</code> , <code>weight_variable</code> , <code>technology</code> and <code>matching</code> variables, and any variables used in <code>disagscors_list</code>).
<code>output_variable</code>	Character scalar. Name of the dependent (output) variable used in the production frontier.
<code>input_variables</code>	Character vector of input (production) variables entering the stochastic frontier.
<code>inefficiency_covariates</code>	Optional named list specifying variables in the inefficiency function. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> .
<code>risk_covariates</code>	Optional named list specifying variables in the production risk (noise) function. If <code>NULL</code> , a homoskedastic noise term is usually implied.
<code>weight_variable</code>	Optional character scalar giving the name of the sampling weight variable in <code>data</code> . Used when computing weighted summaries (e.g., weighted means) of efficiency scores.
<code>production_slope_shifters</code>	Character scalar naming a variable that shifts the production-function slope (e.g., technology shifter). Defaults to <code>NULL</code> to indicate no slope shifter.
<code>intercept_shifters</code>	Optional named list of intercept shifter variables for the baseline (unmatched) sample. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> .
<code>f</code>	Functional form identifier passed to <code>msf_workhorse()</code> (e.g., Cobb-Douglas, translog). The exact meaning is handled by the workhorse function.
<code>d</code>	Distributional form identifier for the inefficiency term (e.g., half-normal, truncated-normal), passed through to <code>msf_workhorse()</code> .
<code>identifiers</code>	Character vector of variable names that uniquely identify units (e.g., household, plot, or observation IDs). These are used to merge efficiency scores back to <code>data</code> and for disaggregated summaries.
<code>include_trend</code>	Logical; if <code>TRUE</code> , includes a technology trend variable (given by <code>technology_variable</code>) in the frontier. Defaults to <code>FALSE</code> .
<code>technology_variable</code>	Optional character scalar naming the technology (trend) variable to be used when <code>include_trend = TRUE</code> .
<code>matching_type</code>	Optional variable or object controlling nearest-neighbor matching; passed to <code>msf_workhorse()</code> . The expected type and structure depend on the implementation of that workhorse function.
<code>adoption_covariates</code>	Optional named list specifying inefficiency-function variables for the matched sample (post-matching specification). Same structure as <code>inefficiency_covariates</code> .
<code>intercept_shifters_meta</code>	Optional named list of intercept shifters for the matched sample. Same structure as <code>intercept_shifters</code> .

```

disagscors_list
  Optional character vector of variable names for which disaggregated efficiency
  score summaries should be computed. For each variable in this list, the function
  builds weighted and unweighted summaries of TE/TE0/TGR/MTE by survey,
  sample, and disaggregation level.

match_specifications
  match specifications
match_specification_optimal
  match specifications
match_path    match path

```

Details

The main outputs include model estimates, mean/percentile summaries, and (optionally) disaggregated efficiency statistics by user-specified grouping variables. For draws other than zero, sample-level results are dropped to reduce storage.

- 1. Draw-specific filtering:** Using `drawlist`, the function removes any `EaId` values associated with the current draw from `data`. This allows for bootstrap, jackknife, or matched-sample resampling.
- 2. Survey-specific estimation:** After setting the `Surveyy` label (either from `Surveyx` or collapsed to "GLSS0"), the function loops over unique survey labels and calls `msf_workhorse()` for each. The workhorse is expected to return a list with components such as `sf_estm`, `el_mean`, `ef_mean`, `rk_mean`, `ef_dist`, `rk_dist`, `el_samp`, `ef_samp`, and `rk_samp`.
- 3. Disaggregated efficiency summaries:** If `disagscors_list` is provided, the function constructs weighted and unweighted efficiency statistics (e.g., weighted mean, mean, median, mode) for TE/TE0/TGR/MTE by survey, sample, restriction status, technology, and disaggregation level. Results are stored in `res$disagscors`.
- 4. Draw-specific pruning:** When `draw != 0`, sample-level objects (`el_samp`, `ef_samp`, `rk_samp`) are removed from the returned list to reduce memory usage, keeping only aggregate results.

Internally, the function uses **data.table**, **dplyr**, **tidyr**, **doBy**, and **crayon** for data manipulation and progress messages. All estimation is delegated to `msf_workhorse()`.

Value

A named list with components:

- `sf_estm` - Combined frontier parameter estimates across surveys (rows tagged with `Surveyy` and `draw`).
- `el_mean`, `ef_mean`, `rk_mean` - Mean/summary statistics for elasticities, efficiency, and risk metrics.
- `ef_dist`, `rk_dist` - Distributions of efficiency and risk quantities.
- `el_samp`, `ef_samp`, `rk_samp` - Sample-level results (only retained when `draw == 0`).
- `disagscors` - Optional disaggregated efficiency summary table if `disagscors_list` is not `NULL`; otherwise `NULL`.

If an error occurs anywhere in the pipeline, the function returns `NULL`.

`draw_msf_summary` *Summarize multi-draw Meta Stochastic Frontier (MSF) results*

Description

Aggregates results from multiple stochastic frontier analysis (SFA) / MSF draws (e.g., from a bootstrap or jackknife procedure) and computes jackknife-style summary statistics for coefficients, efficiency scores, elasticities, risk measures, and their distributions.

Usage

```
draw_msf_summary(res, technology_legend)
```

Arguments

<code>res</code>	List of draw-level result objects. Each element is expected to be a list with (at least) the following components:
	<ul style="list-style-type: none"> • <code>sf_estm</code>: coefficient-level estimates and diagnostics for naive, group, and meta-frontiers, including a <code>draw</code> index. • <code>ef_mean</code>: aggregated efficiency statistics (e.g., TE0, TE, TGR, MTE) by sample, technology, survey, and estimation type. • <code>el_mean</code>: aggregated elasticity statistics by input, technology, and survey. • <code>rk_mean</code>: aggregated risk statistics, if available. • <code>ef_dist</code>: histogram-style efficiency distributions (counts/weights) over value ranges (<code>binned_range_name</code>, <code>binned_range_level</code>). • <code>rk_dist</code>: histogram-style risk distributions, if available. • <code>disagscors</code>: optional disaggregated efficiency summaries by subgroup levels. • <code>el_samp</code>, <code>ef_samp</code>, <code>rk_samp</code>: observation-level elasticities, efficiencies, and risk measures for at least the baseline draw (usually <code>draw == 0</code>).
	Each component must include a <code>draw</code> column that distinguishes the baseline draw (typically <code>draw == 0</code>) from resampled draws (<code>draw != 0</code>).
<code>technology_legend</code>	Data.frame providing the mapping between numeric technology codes and their labels, typically with at least: <ul style="list-style-type: none"> • <code>Tech</code>: numeric technology index used internally in MSF estimation, and • a corresponding label column (e.g., the original <code>technology_variable</code>).

This is used to compute technology gaps (e.g., efficiency or disaggregated efficiency gaps) relative to the reference technology (smallest value in `technology_legend$Tech`).

Details

The function assumes each element of `res` is the output of a single draw from an MSF pipeline (e.g., `msf_workhorse()` followed by `draw_msf_estimations()`), containing components such as: `sf_estm`, `ef_mean`, `el_mean`, `rk_mean`, `ef_dist`, `rk_dist`, and (optionally) `disagscors`.

For each component, the function:

1. Stacks all draws using `data.table::rbindlist()`.

2. Drops non-finite values (NA, Inf, -Inf, NaN).
3. For each statistic (e.g., coefficient estimate, mean efficiency), identifies the baseline draw (`draw == 0`) and joins it with resampled draws (`draw != 0`) aggregated via `doBy::summaryBy()` to compute:
 - `Estimate.mean`: mean across non-baseline draws,
 - `Estimate.sd`: standard deviation across non-baseline draws,
 - `Estimate.length`: number of non-baseline draws.
4. Computes a jackknife-style test statistic and p-value:
 - `jack_zv = Estimate / Estimate.sd`,
 - `jack_pv = 2 * (1 - pt(|jack_zv|, df = Estimate.length))`,

where `Estimate` is the baseline draw statistic and the distribution of resampled draws is treated as a reference distribution.

For disaggregated scores (`disagscors`), the function also:

- Computes technology-specific disaggregated efficiency gaps (level and percent) relative to the reference technology in `technology_legend`, and
- Aggregates these gaps across draws in the same jackknife fashion as for the main efficiency statistics.

Value

A list with the following components:

- `sf_estm`: Data.frame of jackknife summary statistics for frontier coefficients and diagnostics across draws, with columns such as `Estimate`, `Estimate.mean`, `Estimate.sd`, `Estimate.length`, `jack_zv`, and `jack_pv`, indexed by `Survey`, `CoefName`, `Tech`, `sample`, and `restrict`.
- `ef_mean`: Jackknife summaries for efficiency statistics (TE0, TE, TGR, MTE), by sample, technology, survey, type, estimation type, statistic, and restriction.
- `el_mean`: Jackknife summaries for elasticity statistics, by input, technology, survey, statistic, and restriction.
- `rk_mean`: (If available) jackknife summaries for risk statistics, with `type = "risk"`.
- `ef_dist`: Jackknife summaries of efficiency distributions (counts and weights) over ranges, with `stat` indicating whether the row refers to counts or weights.
- `rk_dist`: (If available) jackknife summaries of risk distributions over ranges, with `type = "risk"`.
- `disagscors`: (If available) jackknife summaries for disaggregated efficiency statistics and their gaps by subgroup level and technology.
- `el_samp`: Observation-level elasticity outputs from the first draw in `res` (typically the baseline draw).
- `ef_samp`: Observation-level efficiency scores from the first draw in `res`.
- `rk_samp`: Observation-level risk measures from the first draw in `res`.

`get_crop_area_list` *Extract Matching Crop Area Columns from a Dataset*

Description

Identifies and returns the column names in a dataset corresponding to crop area variables for a specified set of crops. The function looks for column names that start with "Area_" and match any of the crops provided in `selected_crops`.

Usage

```
get_crop_area_list(
  data,
  selected_crops = c("Beans", "Cassava", "Cocoa", "Cocoyam", "Maize", "Millet",
    "Palm", "Peanut", "Pepper", "Plantain", "Rice", "Sorghum", "Tomatoe", "Yam")
)
```

Arguments

<code>data</code>	A <code>data.frame</code> or <code>data.table</code> containing crop-related variables. Column names are expected to include fields prefixed with "Area_", such as "Area_Maize".
<code>selected_crops</code>	A character vector specifying the crop names to filter. Defaults to a common set of crops including "Beans", "Cassava", "Cocoa", "Cocoyam", "Maize", "Millet", "Okra", "Palm", "Peanut", "Pepper", "Plantain", "Rice", "Sorghum", "Tomatoe", and "Yam".

Value

A character vector containing the names of columns in `data` that correspond to the specified crop area variables (e.g., "Area_Maize", "Area_Rice"). Returns an empty vector if no matching columns are found.

See Also

Other helpers: `compute_jenks_binned_shares()`, `harmonized_data_prep()`

`get_household_data` *Download and Load Harmonized Household Data from the GHAGricProductivityLab Repository*

Description

Retrieves a harmonized household- or farm-level dataset from the **GHAGricProductivityLab** GitHub repository using **piggyback**, stores it in a package-specific cache directory, and returns the dataset as a `data.frame`. The file is downloaded only once and reused from the local cache on future calls.

Usage

```
get_household_data(
  dataset = "harmonized_crop_farmer_data",
  github_token = NULL
)
```

Arguments

dataset	Character string. Base name of the dataset to retrieve (without file extension). Must correspond to a .dta file in the hh_data GitHub release (e.g., "harmonized_crop_farm")
github_token	Optional GitHub personal access token (PAT). If NULL, the function checks the environment variable GHProdLab_TOKEN. If that is also missing, the piggy-back download will use default authentication behavior.

Details

The function downloads a Stata .dta file associated with the chosen dataset from the GitHub release labeled hh_data. It uses the package-specific cache directory determined by:

```
tools::R_user_dir("GHAgricProductivityLab", which = "cache")
```

If the file already exists locally, it is not re-downloaded.

GitHub Authentication

- If github_token is supplied, it is used.
- Otherwise, the function looks for environment variable GHProdLab_TOKEN.
- If neither is available, the function falls back to default GitHub credentials (e.g., from gh CLI or cached credentials).

Value

A data.frame containing the requested harmonized dataset.

harmonized_data_prep

Prepare Data for Agricultural Productivity Analysis

Description

Cleans and transforms a dataset by creating new variables, applying log transformations, converting selected variables to factors or characters, and recoding education levels. The function is designed to standardize inputs for further analysis of agricultural productivity.

Usage

```
harmonized_data_prep(data)
```

Arguments

data	A data.frame or data.table containing household- and farm-level variables such as weights, demographic information, and agricultural inputs.
------	--

Value

A cleaned and transformed `data.frame` or `data.table` with additional variables ready for analysis.

See Also

Other helpers: `compute_jenks_binned_shares()`, `get_crop_area_list()`

`match_sample_specifications`

Draw / Match Sample Specifications

Description

Generates (i) a draw list by sampling `EaId` within each unique `Survey` group and (ii) a grid of matching specifications for each bootstrap draw.

Usage

```
match_sample_specifications(number_of_draws = NULL, data, myseed = NULL)
```

Arguments

<code>number_of_draws</code>	Integer. The number of draws to perform per <code>Survey</code> .
<code>data</code>	A <code>data.frame</code> or <code>data.table</code> containing at least the columns <code>Survey</code> and <code>EaId</code> .
<code>myseed</code>	Integer. Seed for random number generation (default 03242025).

Details

Draw list: For each unique value of `Survey`, the function samples `number_of_draws` `EaId` values with replacement and prepends a 0 row (ID = 0) for the baseline. The result is then spread to wide format with one column per `Survey`.

Matching specs: For each draw ID, creates a set of matching model specifications that include:

- Nearest neighbor with distances: `"euclidean"`, `"scaled_euclidean"`, `"mahalanobis"`, `"robust_mahalanobis"`.
- Nearest neighbor with distance `"glm"` and links: `"logit"`, `"probit"`, `"cloglog"`, `"cauchit"`.

An `ARRAY` index is added for convenience.

Value

A list with two elements:

`m.specs` A `data.frame` of matching specifications with columns `boot`, `method`, `distance`, `link`, `ARRAY`.

`drawlist` A `data.frame` in wide format where each `Survey` is a column and rows correspond to draw ID (0:`number_of_draws`).

Note

Requires that `data` contain `Survey` and `Entity`. `tidyR::spread()` is used for wide reshaping (consider `tidyR::pivot_wider()` in new code).

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `treatment_effect_calculati`
`treatment_effect_summary()`, `write_match_formulas()`

`msf_workhorse`

Meta stochastic frontier (MSF) workhorse

Description

Performs Meta Stochastic Frontier (MSF) analysis in several stages:

1. **Naive SF:** Fits a pooled stochastic frontier and computes baseline technical efficiency (TE) measures.
2. **Group SF:** If a technology variable `technology_variable` is supplied, splits the sample into technology groups and fits group-specific SF models.
3. **Meta SF (unmatched and matched):** Constructs a meta-frontier using fitted values from group models and, optionally, nearest-neighbor matching specifications. From this, it derives technology gap ratios (TGR) and meta-technical efficiency (MTE).
4. **Summaries and distributions:** Produces weighted and unweighted summaries and distributional statistics for efficiency, elasticity, and risk across technologies, samples (unmatched/matched), and surveys.

Usage

```
msf_workhorse(
  data,
  output_variable,
  input_variables,
  inefficiency_covariates = NULL,
  risk_covariates = NULL,
  weight_variable = NULL,
  production_slope_shifters = NULL,
  intercept_shifters = NULL,
  f,
  d,
  identifiers,
  include_trend = FALSE,
  technology_variable = NULL,
  matching_type = NULL,
  adoption_covariates = NULL,
  intercept_shifters_meta = NULL,
  match_path,
  match_specifications,
  match_specification_optimal
)
```

Arguments

<code>data</code>	A data.frame or data.table containing the estimation sample. Must include the dependent variable <code>output_variable</code> , the inputs in <code>input_variables</code> , the weight variable <code>weight_variable</code> , unique ID variables in <code>identifiers</code> , any inefficiency and risk covariates in <code>inefficiency_covariates</code> and <code>risk_covariates</code> , the technology variable <code>technology_variable</code> (if used), and any variables required by matching objects on disk (e.g., "Surveyx", "EaId", "HhId", "Mid", "unique_identifier").
<code>output_variable</code>	Character scalar. Name of the dependent (output) variable used in the production frontier.
<code>input_variables</code>	Character vector of input variable names (e.g., land, labor, capital) passed through to <code>sf_workhorse</code> .
<code>inefficiency_covariates</code>	Optional named list specifying variables in the inefficiency function for the naive and group frontiers. Typical structure: <code>list(scalar_variables = c(...), factor_variables = c(...))</code> .
<code>risk_covariates</code>	Optional named list specifying variables in the production risk (noise) function for the naive and group frontiers. Same structure as <code>inefficiency_covariates</code> . If <code>NULL</code> , a homoskedastic noise term is usually implied.
<code>weight_variable</code>	Character scalar. Name of the sampling/observation weight variable in <code>data</code> . Observations with zero weight are dropped.
<code>production_slope_shifters</code>	Character scalar giving the name of a slope-shifter variable in <code>data</code> (e.g., technology shifter, policy dummy). Defaults to <code>NULL</code> for no slope shifter.
<code>intercept_shifters</code>	Optional named list of intercept shifter variables for the naive and group frontiers. Typical structure: <code>list(scalar_variables = c(...), factor_variables = c(...))</code> , passed to <code>sf_workhorse</code> .
<code>f</code>	Functional form index for the production frontier (e.g., Cobb-Douglas, translog, quadratic). The index is passed to <code>sf_workhorse</code> and ultimately to the internal form-selection utilities (e.g., <code>sf_functional_forms</code>).
<code>d</code>	Distributional form index for the inefficiency term (e.g., half-normal, truncated-normal). Passed to <code>sf_workhorse</code> .
<code>identifiers</code>	Character vector of variable names that uniquely identify observations (e.g., <code>c("unique_identifier", "Survey", "CropID", "HhId", "EaId", "Mid")</code>). These IDs are used when merging scores and summarizing by technology or sample.
<code>include_trend</code>	Logical; if <code>TRUE</code> , the last element in <code>input_variables</code> is treated as a trend/technology variable in the production function. Passed through to <code>sf_workhorse</code> . Defaults to <code>FALSE</code> .
<code>technology_variable</code>	Optional character scalar naming the technology variable (e.g., region, system, period) used to define groups for group SF and meta-frontier estimation. If <code>NULL</code> , no technology groups are formed and only the naive frontier and TE are computed.

```

matching_type
    Optional character scalar controlling which nearest-neighbor matching specifications to use for meta-frontier estimation. When not NULL, the function reads matching specifications and matched samples from "results/matching/" and related RDS files. Setting matching_type = "optimal" restricts attention to the subset of matching specifications labeled as optimal.

adoption_covariates
    Optional named list specifying inefficiency-function covariates for the meta-frontier (matched/unmatched) estimation. If NULL, defaults to inefficiency_covariates.

intercept_shifters_meta
    Optional named list of intercept shifters for the meta-frontier estimation. If NULL, defaults to intercept_shifters.

match_path      match path
match_specifications
    match specifications
match_specification_optimal
    match specifications

```

Details

The function is designed as a high-level orchestrator around `sf_workhorse`, assembling naive, group-level, and meta-frontier results into a unified output structure suitable for MSF analysis.

The workflow can be summarized as:

- **Naive SFA (TE0):** Calls `sf_workhorse` on the full sample to obtain baseline efficiencies (teBC, teJLMS, teMO). If no technology variable is specified (`technology_variable` = NULL), these naive efficiencies are the final scores and only TE-related summaries are produced.
- **Group SFA (TE):** When `technology_variable` is provided, the sample is recoded into numeric technology groups (Tech), and `sf_workhorse` is applied separately to each group. Group-level efficiencies are combined into TE measures by group and survey.
- **Meta-frontier (TGR and MTE):** Using fitted group-frontier values (Yhat) from the restricted models, the function fits meta-frontiers (unmatched and, optionally, matched) again via `sf_workhorse`. Technology gap ratios (TGR) are derived from these meta-frontiers, and meta-technical efficiency (MTE) is computed as TE * TGR.
- **Matching layer (optional):** If `matching_type` is given, the function loads matching specifications (mspecs, mspec_optimal) and corresponding matched samples from disk (e.g., "results/matching/"), fits additional meta-frontiers by matching specification, and augments the set of scores and summaries with these matched samples.
- **Summaries and distributions:**

Scores: TE0, TE, TGR, and MTE are reshaped into long form and aggregated to compute weighted/unweighted means, medians, modes, and distribution histograms (both count-based and weight-based) across surveys, samples, technologies, and estimation types. When `technology_variable` is not NULL, the function also computes technology gaps (level and percent) relative to the minimum technology in `technology_legend`.

Elasticities: Elasticities from the underlying SF models (naive, group, meta) are combined and summarized in a similar way, including technology-gap metrics for elasticities when `technology_variable` is provided.

Risk: Risk measures derived from `sf_workhorse` are combined, summarized, and used to build distributional statistics analogous to those for efficiency.

- **LR tests:** When `technology_variable` is not NULL, the function builds likelihood ratio test statistics comparing a naive pooled frontier to the combination of group and meta-frontiers, storing these as rows with `CoefName = "LRT"` in the `sf_estm` output.

Throughout, the function uses `dplyr`, `tidyr`, `data.table`, `doBy`, and related helper functions for reshaping and summarizing the output of `sf_workhorse`.

Value

A named list with the following components:

- `sf_estm`: Data.frame of coefficient-level summaries from all naive, group, and meta-frontiers (including LR tests), augmented with technology identifiers (`Tech`) and sample labels (e.g., "unmatched", matching-spec names).
- `ef_mean`: Data.frame of aggregated efficiency statistics (TE0, TE, TGR, MTE), including weighted/unweighted means, medians, and modes, by survey, sample, technology, type, estimation type, and restriction status. When `technology_variable` is supplied, also includes efficiency gap levels and percentages.
- `ef_dist`: Data.frame of efficiency distributions (histogram counts and weights) over efficiency ranges, by survey, sample, technology, type, estimation type, and restriction.
- `rk_dist`: Data.frame of risk distributions (histogram counts and weights) analogous to `ef_dist`, with `type = "risk"`.
- `el_mean`: Data.frame of aggregated elasticity statistics (weighted/unweighted means, medians, modes) by input, survey, sample, technology, and restriction, including elasticity-gap measures when `technology_variable` is provided.
- `rk_mean`: Data.frame of aggregated risk statistics (weighted and unweighted) by survey, sample, technology, and restriction, with risk-gap measures when `technology_variable` is provided.
- `el_samp`: Data.frame of observation-level elasticities for all models and samples (excluding the synthetic "GLSS0" survey rows).
- `ef_samp`: Data.frame of observation-level efficiency scores (TE0, TE, TGR, MTE) including IDs, weights, technologies, and sample labels.
- `rk_samp`: Data.frame of observation-level risk measures for all models and samples (excluding the synthetic "GLSS0" survey rows).

See Also

Other frontier analysis: `sf_workhorse()`

`run_only_for`

Restrict Execution of an R Script to Allowed SLURM Job Conditions

Description

`run_only_for()` controls when an R script is allowed to run based on:

- the SLURM array index (`SLURM_ARRAY_TASK_ID`)
- the SLURM job name (`SLURM_JOB_NAME`)
- whether the script is being run locally on Windows

The function is designed for workflows where multiple R scripts run in a single SLURM array job; each script decides independently whether it should run.

Usage

```
run_only_for(id = NULL, run_on_windows = TRUE, allowed_jobnames = NULL)
```

Arguments

<code>id</code>	Integer or NULL. Expected SLURM array index, or NULL to skip checking (useful for Windows/local runs).
<code>run_on_windows</code>	Logical. When TRUE (default), do not restrict when running on Windows.
<code>allowed_jobnames</code>	Character vector of SLURM job names allowed to run this script. If NULL, job-name filtering is skipped.

Details

Execution logic:

1. **Windows override** If running on Windows and `run_on_windows` = TRUE, return immediately.
2. **Allowed job names** If `allowed_jobnames` is set, the script runs only when SLURM job name is in that vector.
3. **Array ID restriction**
 - If `id` = NULL, skip array-index filtering.
 - If numeric, match against `SLURM_ARRAY_TASK_ID`.

Any violation -> quit (`save = "no"`).

Value

Invisibly returns NULL when allowed; otherwise exits R.

`sf_workhorse` *Stochastic frontier workhorse*

Description

Fits a stochastic frontier model for a given production specification and, when monotonicity is sufficiently violated, re-estimates a shape-constrained (restricted) frontier using minimum-distance methods. The function wraps the underlying sfaR estimation routines and a set helper utilities (e.g., `sf_functional_forms()`, `equation_editor()`, `fit_organizer()`, `translogEla()`, curvature/monotonicity checks) to produce:

1. Unrestricted stochastic frontier estimates and diagnostics.
2. Shape-constrained estimates that enforce monotonicity (and related regularity conditions) when needed.
3. Observation-level efficiencies, fitted values, and risk measures.
4. Observation-level elasticities, including a "returns-to-scale" term.

Usage

```
sf_workhorse(
  data,
  output_variable,
  input_variables,
  weight_variable = NULL,
  d,
  f,
  production_slope_shifters = NULL,
  intercept_shifters = NULL,
  inefficiency_covariates = NULL,
  risk_covariates = NULL,
  identifiers,
  include_trend = FALSE
)
```

Arguments

<code>data</code>	A data.frame or data.table containing all variables required for estimation, including the dependent variable <code>output_variable</code> , inputs <code>input_variables</code> , the weight variable <code>weight_variable</code> (if used), the unique ID variables listed in <code>identifiers</code> , and any variables referenced in <code>production_slope_shifters</code> , <code>intercept_shifters</code> , <code>inefficiency_covariates</code> , or <code>risk_covariates</code> .
<code>output_variable</code>	Character scalar. Name of the dependent/output variable in <code>data</code> used for the production frontier.
<code>input_variables</code>	Character vector of input variable names (e.g., land, labor, capital). The order of variables in <code>input_variables</code> determines how they map into the generic input labels <code>I1</code> , <code>I2</code> , ... used in the functional-form utilities.
<code>weight_variable</code>	Optional character scalar. Name of the sampling/observation weight variable. If <code>NULL</code> , all observations are given unit weight.
<code>d</code>	Distributional form index for the inefficiency term. This is passed to <code>sf_functional_forms()</code> and ultimately to <code>sfacross()</code> (e.g., to select half-normal, truncated-normal, etc.). The exact mapping is determined by <code>sf_functional_forms()</code> .
<code>f</code>	Functional form index for the production frontier (e.g., Cobb-Douglas, translog, quadratic), used to select from the list returned by <code>sf_functional_forms()</code> . The name of the chosen form (e.g., "CD", "TL", "QD", "LN") influences logging of variables and regularity checks.
<code>production_slope_shifters</code>	Character scalar giving the name of a slope-shifter variable in <code>data</code> (e.g., technology shift, policy dummy). Defaults to <code>NULL</code> for no slope shifter and is passed to <code>equation_editor()</code> .
<code>intercept_shifters</code>	Optional named list of intercept shifter variables for the production function. Typical structure: <code>list(scalar_variables = c(...), factor_variables = c(...))</code> , where the specific interpretation is handled inside <code>equation_editor()</code> .
<code>inefficiency_covariates</code>	Optional named list describing the inefficiency-function covariates. Typical structure: <code>list(scalar_variables = c(...), factor_variables = c(...))</code> , passed as <code>uhet</code> and <code>muhet</code> to <code>sfacross()</code> when non- <code>NULL</code> .

<code>risk_covariates</code>	Optional named list describing the production-risk (noise) covariates. When non-NULL, used as <code>vhet</code> in <code>sfacross()</code> to allow heteroskedasticity of the noise term.
<code>identifiers</code>	Character vector of variable names that uniquely identify observations (e.g., <code>c("unique_identifier", "Survey", "CropID", "HhId", "EaId", "Mid")</code>). These IDs are carried through to efficiency, elasticity, and risk outputs.
<code>include_trend</code>	Logical; if TRUE, the last element in <code>input_variables</code> is treated as a trend/technology variable and handled accordingly in the functional-form and elasticity calculations. Defaults to FALSE.

Details

The function proceeds in several steps:

1. **Setup and functional form selection:** Using `sf_functional_forms()`, the function determines the appropriate production functional form and distributional specification based on the number of inputs (`number_of_inputs`), the `f` and `d` indices, and the `include_trend` flag. This includes whether the model is specified in logs (e.g., CD/TL/GP/TP) or levels.
2. **Variable construction and equation building:** Inputs in `input_variables` are mapped to generic labels (`I1, I2, ...`) and their log transforms are created when needed. The outcome variable is set to `Y` or `lnY`. The function then calls `equation_editor()` to construct the production, inefficiency, and risk equations used by `sfacross()`.
3. **Unrestricted stochastic frontier estimation:** The function iterates over a grid of optimization methods (`nr, nm, bfgs, bhhh, ...`) and tolerance settings, calling `sfacross()` until a successful convergence is reported. It extracts observation-level efficiencies via `sfaR::efficiencies()`, constructs fitted values, and builds an observation-level risk metric based on squared deviations from mean output.
4. **Elasticities and regularity checks:** Using `fit_organizer()` and `translogEla()`, the function computes elasticities and returns-to-scale-type measures, and evaluates monotonicity and curvature via `translogCheckMono()` and `translogCheckCurvature()` (or a coefficient-sign check for CD/LN forms).
5. **Shape-constrained frontier (if needed):** If the monotonicity measure `mono` falls below 0.80, the function constructs a set of linear restrictions using `translogMonoRestr()` and solves a quadratic programming problem (via `quadprog::solve.QP`) and fall-back matrix inversions using **Matrix**, **MASS**, **corpcor**) to obtain constrained coefficients. For TL/QD forms a constrained frontier (`lcFitted`) is computed with `translogCalc()`; for CD/LN forms, constrained coefficients are obtained via a SEM representation using **lavaan**.
6. **Re-estimation under constraints:** Given the constrained frontier, the function rebuilds the production equation with `equation_editor()` and re-estimates a stochastic frontier (`sfc`) using the same optimization grid. Constrained efficiencies, risks, and elasticities are computed and summarized, and regularity checks are repeated.

Unrestricted and restricted summaries are then stacked, with a `restrict` flag ("Unrestricted" vs. "Restricted"), for downstream comparison.

Value

A named list with four data.frames:

- `sf`: Coefficient-level results combining unrestricted and, when applicable, restricted SFA estimates, plus monotonicity/curvature diagnostics. Columns include (at minimum) `CoefName`, `Estimate`, `StdError`, `Zvalue`, `Pvalue`, and `restrict`.
- `ef`: Observation-level efficiency results (unrestricted and restricted), including the ID variables in `identifiers`, `weights`, efficiency measures (e.g., `u`), model-fitted values (`mlFitted`), and `restrict`.
- `el`: Observation-level elasticities and returns-to-scale-type measures. Includes IDs, weights, elasticity columns (e.g., `el1`, `el2`, ...), and the summed elasticity), and `restrict`.
- `rk`: Observation-level risk measures (unrestricted and restricted), including IDs, weights, and `risk`, plus `restrict`.

See Also

Other frontier analysis: [msf_workhorse\(\)](#)

`study_setup`

Initialize Study Environment and Directory Structure

Description

Initialize Study Environment and Directory Structure

Usage

```
study_setup(
  myseed = 1980632,
  project_name,
  local_directories = list(home = file.path("replications", project_name), output =
    file.path("replications", project_name, "output"), matching =
    file.path("replications", project_name, "output", "matching"), treatment_eff =
    file.path("replications", project_name, "output", "treatment_effects"), esti =
    file.path("replications", project_name, "output", "estimations"), figure_dat =
    file.path("replications", project_name, "output", "figure_data"), figure =
    file.path("replications", project_name, "output", "figure")))
)
```

Arguments

<code>myseed</code>	Numeric/integer scalar seed. Will be coerced to integer. Default: 1980632.
<code>project_name</code>	Length-1, non-NA character project name (required).
<code>local_directories</code>	Named list of character paths to create. Defaults use <code>project_name</code> .

Value

List with `wd` (directories) and `seed`.

`treatment_effect_calculation`
Compute log-linear treatment effects (ATE/ATET/ATEU) for multiple outcomes

Description

For a given matching specification index `i`, this function:

1. Loads the corresponding matched weights file from `matching_output_directory`,
2. Merges those weights into the analysis data via `unique_identifier`,
3. Optionally normalizes each outcome by the first element of `outcome_variables` (e.g., to construct per-hectare or per-unit measures),
4. Fits a weighted log-linear model for each outcome with interactions between `Treat` and the supplied matching covariate formulas,
5. Transforms fitted differences into percentage treatment effects, trims extreme values, and computes weighted ATE, ATET, and ATEU alongside model fit statistics.

Usage

```
treatment_effect_calculation(
  data,
  outcome_variables,
  normalize = NULL,
  i,
  match_specifications,
  matching_output_directory,
  match_formulas
)
```

Arguments

<code>data</code>	A <code>data.frame</code> containing at least: <ul style="list-style-type: none"> • <code>unique_identifier</code>: unit identifier used for merging matched weights, • <code>Treat</code>: treatment indicator (logical or 0/1), • <code>Area</code>: plot size (positive; used for input checks and typical scaling), • all variables named in <code>outcome_variables</code>, • any covariates referenced in <code>match_formulas\$general_match</code> and <code>match_formulas\$exact_match</code>.
<code>outcome_variables</code>	Character vector of outcome variable names (e.g., <code>c("HrvstKg", "Area", "SeedKg", "HHLab")</code>)
<code>normalize</code>	Logical scalar. If TRUE, each outcome in <code>outcome_variables</code> is divided by the first element of <code>outcome_variables</code> in the merged data (e.g., to obtain outcomes per hectare or per unit).
<code>i</code>	Integer index selecting the row of <code>match_specifications</code> and the corresponding matching result file (zero-padded via <code>ARRAY</code> to "NNNN.rds").
<code>match_specifications</code>	A <code>data.frame</code> with at least a column <code>ARRAY</code> . The <code>i</code> -th row is used to locate the matching output file and is <code>cbind</code> -ed to the returned results.

```

matching_output_directory
  Directory containing per-specification matching results named "match_0001.rds",
  "match_0002.rds", etc. Each RDS is expected to contain an element $md
  with columns unique_identifier and weights.

match_formulas
  A list with:
    • general_match: a formula; the RHS is accessed via as.character(...)[3],
    • exact_match: a formula; the RHS is accessed via as.character(...)[2].

```

Details

For each outcome `oc` in `outcome_variables`, the function:

1. Filters to observations with non-missing `Treat`, the outcome, and `weights`,
2. Constructs a log-linear model of the form

`log(oc + eps) ~ Treat * (general_match_rhs + exact_match_rhs)`

where `eps = min(oc[oc > 0]) * 1/100` in the estimation sample,

3. Fits the model using `lm()` with weights given by the matched `weights`,
4. Predicts fitted log outcomes for treated (`Treat = 1`) and untreated (`Treat = 0`) counterfactuals,
5. Computes individual percentage treatment effects

$$TE_{OLS} = \{\exp(\hat{y}_1 - \hat{y}_0) - 1\} \times 100,$$

trims them to the interval $[-100, 100]$, and then forms:

- ATE: weighted mean of `TE_OLS` over all units,
- ATET: weighted mean of `TE_OLS` among treated units,
- ATEU: weighted mean of `TE_OLS` among untreated units.

In addition, for each outcome the function records model diagnostics: `AIC`, log-likelihood, R^2 , adjusted R^2 , F-statistic, and the sample size used in the regression.

Value

A `data.frame` (`atet_scalar`) with the columns from `match_specifications[i,]` repeated across rows, and the additional columns:

- `outcome`: outcome variable name,
- `treatment`: name of the treatment indicator (always "`Treat`"),
- `level`: one of "ATE", "ATET", "ATEU", "aic", "ll", "R2", "N", "Ft", "R2a",
- `est`: numeric estimate corresponding to `level`.

One block of rows is returned per outcome in `outcome_variables`. If a model fails for a given outcome, that outcome is skipped.

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `match_sample_specification`, `treatment_effect_summary()`, `write_match_formulas()`

treatment_effect_summary*Summarize Treatment Effect Estimates Across Matching Specifications***Description**

Aggregates and summarizes treatment effect results stored as .rds files in a specified output directory. The function reads each RDS file, combines all results into a single data frame, filters out invalid estimates, and computes jackknife-style summary statistics (mean, standard error, sample size, t-value, and p-value) by matching specification and outcome.

Usage

```
treatment_effect_summary(treatment_effects_output_directory)
```

Arguments

`treatment_effects_output_directory`

Character string giving the path to the directory containing .rds files with treatment effect results (e.g., outputs of `treatment_effect_calculation()` saved across specifications).

Details

The function proceeds through the following steps:

1. Lists and reads all .rds files in `treatment_effects_output_directory`.
2. Combines them into a single long-format data frame using `data.table::rbindlist()`.
3. Removes rows with invalid estimates (NA, NaN, Inf, or -Inf).
4. Separates base (non-bootstrap) results where `boot == 0` and joins these with jackknife summary statistics computed via `doBy::summaryBy()` across method, distance, link, outcome, and level.
5. Renames summary columns to:
 - `jack_mean`: mean estimate
 - `jack_se`: standard error
 - `jack_n`: sample size (number of replicates)
 - `jack_tvl`: t-value (`jack_mean / jack_se`)
 - `jack_pvl`: two-sided p-value

Value

A `data.frame` containing summarized treatment effect estimates with columns:

- `method`, `distance`, `link`, `outcome`, `level`
- `est` (original estimate for non-bootstrap sample)
- `jack_mean`, `jack_se`, `jack_n`, `jack_tvl`, `jack_pvl`

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `match_sample_specification()`, `treatment_effect_calculation()`, `write_match_formulas()`

```
write_match_formulas
```

Construct Matching Formulas for Exact and General Matching

Description

Builds two types of matching formulas commonly used in propensity score or covariate matching workflows:

- An **exact match** formula specifying categorical variables to match exactly.
- A **general match** formula specifying numeric and factor covariates for distance-based matching.

The function returns both formulas as character strings that can be used in matching functions such as `MatchIt::matchit()` or `Matching::Match()`.

Usage

```
write_match_formulas(  
  match_variables_exact,  
  match_variables_scaler,  
  match_variables_factor  
)
```

Arguments

`match_variables_exact`

A character vector of variable names to be included in the exact match component (e.g., "gender", "region").

`match_variables_scaler`

A character vector of continuous or scalar variable names to include as covariates in the general match formula (e.g., "age", "income").

`match_variables_factor`

A character vector of factor variable names to be included as dummy-coded categorical covariates in the general match formula.

Value

A named list with two elements:

`exact_match` A string representing the exact match formula (factor variables only).

`general_match` A string representing the general match formula, including both continuous and factor covariates on the right-hand side of `Treat ~`.

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `match_sample_specification()`, `treatment_effect_calculation()`, `treatment_effect_summary()`

Index

- * **frontier analysis**
 - msf_workhorse, 14
 - sf_workhorse, 18
- * **helpers**
 - compute_jenks_binned_shares,
 - 2
 - get_crop_area_list, 11
 - harmonized_data_prep, 12
- * **matching**
 - covariate_balance, 3
 - draw_matched_samples, 4
 - match_sample_specifications,
 - 13
 - treatment_effect_calculation,
 - 22
 - treatment_effect_summary, 24
 - write_match_formulas, 25
- compute_jenks_binned_shares, 2,
 - 11, 13
- covariate_balance, 3, 5, 14, 23–25
- draw_matched_samples, 4, 4, 14, 23–25
- draw_msf_estimations, 6
- draw_msf_summary, 9
- get_crop_area_list, 3, 11, 13
- get_household_data, 11
- harmonized_data_prep, 3, 11, 12
- match_sample_specifications, 4, 5,
 - 13, 23–25
- msf_workhorse, 14, 21
- run_only_for, 17
- sf_workhorse, 15–17, 18
- study_setup, 21
- treatment_effect_calculation, 4,
 - 5, 14, 22, 24, 25
- treatment_effect_summary, 4, 5, 14,
 - 23, 24, 25
- write_match_formulas, 4, 5, 14, 23, 24,
 - 25