

Package ‘GHAgriсProductivityLab’

November 16, 2025

Type Package

Title Agricultural Productivity in Ghana

Version 0.0.0.9000

Author Francis Tsiboe [aut, cre] (<<https://orcid.org/0000-0001-5984-1072>>)

Maintainer Francis Tsiboe <ftsiboe@hotmail.com>

Contributor -

Reviewer -

Creator Francis Tsiboe

Description Provides tools and datasets for investigating the drivers of agricultural production shortfalls in Ghana.

It compiles research studies that examine farmer-specific and institutional factors to assess whether inefficiencies arise from technical inefficiency, technology gaps, or both. The package offers empirical evidence to inform policy discussions and interventions aimed at improving agricultural productivity, particularly in contexts with limited access to modern technologies.

License GPL-3 + file LICENSE

URL <https://github.com/ftsiboe/GHAgriсProductivityLab>

BugReports <https://github.com/ftsiboe/GHAgriсProductivityLab/issues>

Encoding UTF-8

Roxxygen list(markdown = TRUE)

RoxxygenNote 7.3.3

VignetteBuilder knitr

Depends R (>= 4.1.0)

Imports MatchIt, data.table, haven, stats, sfaR, piggyback

Remotes github::dylan-turner25/rfcip

Suggests dplyr, knitr, crayon, tidyverse, withr, rmarkdown, lavaan, emdbook, cobalt, quadprog, cli, dplyr, doBy, stringr, MASS, Matrix, corpcor, testthat (>= 3.0.0), car, frontier, micEcon, rgenoud

LazyData true

Cite-us If you find it useful, please consider staring the repository and citing the following studies

- Tsiboe, F. and Turner, D. (2025). ``Incorporating buy-up price loss coverage into the United States farm safety net." *Applied Economic Perspectives and Policy*.

- Tsiboe, F., et al. (2025). ``Risk reduction impacts of crop insurance in the United States.'' *Applied Economic Perspectives and Policy*.
- Gaku, S. and Tsiboe, F. (2024). Evaluation of alternative farm safety net program combination strategies. *Agricultural Finance Review*.

Contents

covariate_balance	2
draw_matched_samples	3
draw_msf_estimations	4
draw_msf_summary	7
equation_editor	9
fit_organizer	11
get_crop_area_list	12
get_household_data	13
harmonized_data_prep	14
match_sample_specifications	14
msf_workhorse	15
sfaR_summary	18
sf_functional_forms	20
sf_model_specifications	21
sf_workhorse	23
study_setup	25
treatment_effect_calculation	26
treatment_effect_summary	28
write_match_formulas	29

Index

31

covariate_balance	<i>Compute Covariate Balance Summaries Across Matching Specs</i>
-------------------	--

Description

Compute Covariate Balance Summaries Across Matching Specs

Usage

```
covariate_balance(match_specifications, matching_output_directory)
```

Arguments

match_specifications	data.frame/data.table with at least columns: boot, ARRAY, and the spec fields stored in RDS (e.g., method, distance, link).
matching_output_directory	Directory containing RDS files named as "0001.rds", "0002.rds", etc.

Details

Reads each matching result RDS (expected to contain `m.out` and `match_specifications`), extracts balance via `cobalt::bal.tab`, reshapes, computes a composite balance $rate = mean((Diff - 0)^2, (KS - 0)^2, (V_{Ratio} - 1)^2)$, and averages by ARRAY, method, distance, link, sample.

Value

A list with:

<code>rate</code>	data.frame of mean balance metrics by spec (Adj sample only) with a composite rate.
<code>bal_tab</code>	long-format balance table per covariate/stat/sample/spec.

See Also

Other matching: `draw_matched_samples()`, `match_sample_specifications()`, `treatment_effect_calculation()`, `treatment_effect_summary()`, `write_match_formulas()`

`draw_matched_samples` *Draw matched samples (stratified bootstrap + matching)*

Description

Performs stratified resampling at the Surveyx, EaId level (excluding EaIds listed for the current bootstrap ID) and computes adjusted sampling weights used for matching. Then fits a matching model per `match_specifications[i,]` using **MatchIt**, with exact matching on Emch and distance/link from `match_specifications`.

Usage

```
draw_matched_samples(
  i,
  data,
  match_variables_exact,
  match_variables_scaler,
  match_variables_factor,
  match_specifications,
  sample_draw_list,
  verbose = FALSE
)
```

Arguments

<code>i</code>	Integer index selecting the row of <code>match_specifications</code> to use.
<code>data</code>	A data.frame/data.table containing at least: Surveyx, EaId, HhId, Mid, unique_identifier, Weight, Treat, plus columns named in Emch, Scle, Fixd.
<code>match_variables_exact</code>	A character vector of variable names to be included in the exact match component (e.g., "gender", "region").

<code>match_variables_scaler</code>	A character vector of continuous or scalar variable names to include as covariates in the general match formula (e.g., "age", "income").
<code>match_variables_factor</code>	A character vector of factor variable names to be included as dummy-coded categorical covariates in the general match formula.
<code>match_specifications</code>	Data frame of matching specifications with columns <code>boot</code> , <code>method</code> , <code>distance</code> , and optionally <code>link</code> .
<code>sample_draw_list</code>	Data frame where column ID identifies the bootstrap draw, and each remaining column corresponds to a Survey containing the sampled EaId.
<code>verbose</code>	Logical; if TRUE, prints the chosen matching method and timing. Default FALSE.

Details

Adjusted weights are computed as $pWeight = Weight \times (alloc/allocj)$, where `alloc` is the pre-exclusion sum of `Weight` by Surveyx, `EaId` and `allocj` is the post-exclusion sum.

Exact matching is performed on variables in `Emch`. The distance model formula is constructed as `Treat ~ Scle + Fixd`. When `distance == "glm"`, `m.order = "largest"` is used; otherwise "closest".

Value

A list with:

- `match_specifications` The selected row from `match_specifications`.
- `m.out` The `matchit` object.
- `md` Matched data: Surveyx, `EaId`, `HhId`, `Mid`, `unique_identifier`, `weights`, `pWeight`.
- `df` The analysis data with adjusted weights: Surveyx, `EaId`, `HhId`, `Mid`, `unique_identifier`, `pWeight`.

See Also

Other matching: [covariate_balance\(\)](#), [match_sample_specifications\(\)](#), [treatment_effect_calculation\(\)](#), [treatment_effect_summary\(\)](#), [write_match_formulas\(\)](#)

`draw_msf_estimations` Run a single MSF draw and summarize stochastic frontier results

Description

Performs one iteration of multi-stage stochastic frontier (MSF) estimation for a given draw. The function (i) filters the analysis sample using a draw-specific exclusion list, (ii) calls `msf_workhorse()` to estimate production, inefficiency, and risk components by survey, and (iii) optionally builds disaggregated efficiency score summaries.

Usage

```
draw_msf_estimations(
  draw,
  drawlist,
  surveyy = FALSE,
  data,
  yvar,
  xlist,
  ulist = NULL,
  vlist = NULL,
  wvar = NULL,
  slope_shifter = "NONE",
  intercept_shifters = NULL,
  f,
  d,
  identifiers,
  include_trend = FALSE,
  tvar = NULL,
  matching_type = NULL,
  ulistM = NULL,
  intercept_shiftersM = NULL,
  disagscors_list = NULL
)
```

Arguments

<code>draw</code>	Integer (or numeric coercible to integer). Draw iteration index. Typically 0 is used for the full baseline sample and positive integers for resampled or matched draws.
<code>drawlist</code>	A data.frame or similar object containing exclusion information by draw. Rows correspond to draws, with column ID identifying the draw and subsequent columns containing EaId values to remove from data for that draw.
<code>surveyy</code>	Logical; if TRUE, uses the survey labels stored in data\$Surveyx. If FALSE (default), all observations are assigned to a synthetic survey labeled "GLSS0" for estimation.
<code>data</code>	A data.frame or data.table containing the estimation sample. Must include, at minimum, the columns referenced in other arguments (e.g., <code>yvar</code> , <code>xlist</code> , <code>identifiers</code> , <code>wvar</code> , technology and matching variables, and any variables used in <code>disagscors_list</code>).
<code>yvar</code>	Character scalar. Name of the dependent (output) variable used in the production frontier.
<code>xlist</code>	Character vector of input (production) variables entering the stochastic frontier.
<code>ulist</code>	Optional named list specifying variables in the inefficiency function. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> .
<code>vlist</code>	Optional named list specifying variables in the production risk (noise) function. If NULL, a homoskedastic noise term is usually implied.
<code>wvar</code>	Optional character scalar giving the name of the sampling weight variable in data. Used when computing weighted summaries (e.g., weighted means) of efficiency scores.

<code>slope_shifter</code>	Character scalar naming a variable that shifts the production-function slope (e.g., technology shifter). Defaults to "NONE" to indicate no slope shifter.
<code>intercept_shifters</code>	Optional named list of intercept shifter variables for the baseline (unmatched) sample. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> .
<code>f</code>	Functional form identifier passed to <code>msf_workhorse()</code> (e.g., Cobb-Douglas, translog). The exact meaning is handled by the workhorse function.
<code>d</code>	Distributional form identifier for the inefficiency term (e.g., half-normal, truncated-normal), passed through to <code>msf_workhorse()</code> .
<code>identifiers</code>	Character vector of variable names that uniquely identify units (e.g., household, plot, or observation IDs). These are used to merge efficiency scores back to data and for disaggregated summaries.
<code>include_trend</code>	Logical; if TRUE, includes a technology trend variable (given by <code>tvar</code>) in the frontier. Defaults to FALSE.
<code>tvar</code>	Optional character scalar naming the technology (trend) variable to be used when <code>include_trend</code> = TRUE.
<code>matching_type</code>	Optional variable or object controlling nearest-neighbor matching; passed to <code>msf_workhorse()</code> . The expected type and structure depend on the implementation of that workhorse function.
<code>ulistM</code>	Optional named list specifying inefficiency-function variables for the matched sample (post-matching specification). Same structure as <code>ulist</code> .
<code>intercept_shiftersM</code>	Optional named list of intercept shifters for the matched sample. Same structure as <code>intercept_shifters</code> .
<code>disagscors_list</code>	Optional character vector of variable names for which disaggregated efficiency score summaries should be computed. For each variable in this list, the function builds weighted and unweighted summaries of TE/TE0/TGR/MTE by survey, sample, and disaggregation level.

Details

The main outputs include model estimates, mean/percentile summaries, and (optionally) disaggregated efficiency statistics by user-specified grouping variables. For draws other than zero, sample-level results are dropped to reduce storage.

1. **Draw-specific filtering:** Using `drawlist`, the function removes any `EaId` values associated with the current draw from `data`. This allows for bootstrap, jackknife, or matched-sample resampling.
2. **Survey-specific estimation:** After setting the Surveyy label (either from `Surveyx` or collapsed to "GLSS0"), the function loops over unique survey labels and calls `msf_workhorse()` for each. The workhorse is expected to return a list with components such as `sf_estm`, `el_mean`, `ef_mean`, `rk_mean`, `ef_dist`, `rk_dist`, `el_samp`, `ef_samp`, and `rk_samp`.
3. **Disaggregated efficiency summaries:** If `disagscors_list` is provided, the function constructs weighted and unweighted efficiency statistics (e.g., weighted mean, mean, median, mode) for TE/TE0/TGR/MTE by survey, sample, restriction status, technology, and disaggregation level. Results are stored in `res$disagscors`.
4. **Draw-specific pruning:** When `draw != 0`, sample-level objects (`el_samp`, `ef_samp`, `rk_samp`) are removed from the returned list to reduce memory usage, keeping only aggregate results.

Internally, the function uses **data.table**, **dplyr**, **tidyr**, **doBy**, and **crayon** for data manipulation and progress messages. All estimation is delegated to `msf_workhorse()`.

Value

A named list with components:

- `sf_estm` - Combined frontier parameter estimates across surveys (rows tagged with Survey and draw).
- `el_mean`, `ef_mean`, `rk_mean` - Mean/summary statistics for elasticities, efficiency, and risk metrics.
- `ef_dist`, `rk_dist` - Distributions of efficiency and risk quantities.
- `el_samp`, `ef_samp`, `rk_samp` - Sample-level results (only retained when `draw == 0`).
- `disagscors` - Optional disaggregated efficiency summary table if `disagscors_list` is not NULL; otherwise NULL.

If an error occurs anywhere in the pipeline, the function returns NULL.

`draw_msf_summary`

Summarize multi-draw Meta Stochastic Frontier (MSF) results

Description

Aggregates results from multiple stochastic frontier analysis (SFA) / MSF draws (e.g., from a bootstrap or jackknife procedure) and computes jackknife-style summary statistics for coefficients, efficiency scores, elasticities, risk measures, and their distributions.

Usage

```
draw_msf_summary(res, technology_legend)
```

Arguments

<code>res</code>	List of draw-level result objects. Each element is expected to be a list with (at least) the following components: <ul style="list-style-type: none"> • <code>sf_estm</code>: coefficient-level estimates and diagnostics for naive, group, and meta-frontiers, including a draw index. • <code>ef_mean</code>: aggregated efficiency statistics (e.g., TE0, TE, TGR, MTE) by sample, technology, survey, and estimation type. • <code>el_mean</code>: aggregated elasticity statistics by input, technology, and survey. • <code>rk_mean</code>: aggregated risk statistics, if available. • <code>ef_dist</code>: histogram-style efficiency distributions (counts/weights) over value ranges (range, Frqllevel). • <code>rk_dist</code>: histogram-style risk distributions, if available. • <code>disagscors</code>: optional disaggregated efficiency summaries by subgroup levels. • <code>el_samp</code>, <code>ef_samp</code>, <code>rk_samp</code>: observation-level elasticities, efficiencies, and risk measures for at least the baseline draw (usually <code>draw == 0</code>).
<code>technology_legend</code>	Each component must include a <code>draw</code> column that distinguishes the baseline draw (typically <code>draw == 0</code>) from resampled draws (<code>draw != 0</code>).

`technology_legend`

Data.frame providing the mapping between numeric technology codes and their labels, typically with at least:

- Tech: numeric technology index used internally in MSF estimation, and
- a corresponding label column (e.g., the original tvar).

This is used to compute technology gaps (e.g., efficiency or disaggregated efficiency gaps) relative to the reference technology (smallest value in `technology_legend$Tech`).

Details

The function assumes each element of `res` is the output of a single draw from an MSF pipeline (e.g., `msf_workhorse()` followed by `draw_msf_estimations()`), containing components such as: `sf_estm`, `ef_mean`, `el_mean`, `rk_mean`, `ef_dist`, `rk_dist`, and (optionally) `disagscors`.

For each component, the function:

1. Stacks all draws using `data.table::rbindlist()`.
2. Drops non-finite values (NA, Inf, -Inf, NaN).
3. For each statistic (e.g., coefficient estimate, mean efficiency), identifies the baseline draw (`draw == 0`) and joins it with resampled draws (`draw != 0`) aggregated via `doBy::summaryBy()` to compute:
 - `Estimate.mean`: mean across non-baseline draws,
 - `Estimate.sd`: standard deviation across non-baseline draws,
 - `Estimate.length`: number of non-baseline draws.
4. Computes a jackknife-style test statistic and p-value:
 - `jack_zv` = `Estimate / Estimate.sd`,
 - `jack_pv` = $2 * (1 - pt(|jack_zv|, df = Estimate.length))$,

where `Estimate` is the baseline draw statistic and the distribution of resampled draws is treated as a reference distribution.

For disaggregated scores (`disagscors`), the function also:

- Computes technology-specific disaggregated efficiency gaps (level and percent) relative to the reference technology in `technology_legend`, and
- Aggregates these gaps across draws in the same jackknife fashion as for the main efficiency statistics.

Value

A list with the following components:

- `sf_estm`: Data.frame of jackknife summary statistics for frontier coefficients and diagnostics across draws, with columns such as `Estimate`, `Estimate.mean`, `Estimate.sd`, `Estimate.length`, `jack_zv`, and `jack_pv`, indexed by `Survey`, `CoefName`, `Tech`, `sample`, and `restrict`.
- `ef_mean`: Jackknife summaries for efficiency statistics (TE0, TE, TGR, MTE), by sample, technology, survey, type, estimation type, statistic, and restriction.
- `el_mean`: Jackknife summaries for elasticity statistics, by input, technology, survey, statistic, and restriction.
- `rk_mean`: (If available) jackknife summaries for risk statistics, with `type = "risk"`.
- `ef_dist`: Jackknife summaries of efficiency distributions (counts and weights) over ranges, with `stat` indicating whether the row refers to counts or weights.
- `rk_dist`: (If available) jackknife summaries of risk distributions over ranges, with `type = "risk"`.

- disagscors: (If available) jackknife summaries for disaggregated efficiency statistics and their gaps by subgroup level and technology.
- el_samp: Observation-level elasticity outputs from the first draw in res (typically the baseline draw).
- ef_samp: Observation-level efficiency scores from the first draw in res.
- rk_samp: Observation-level risk measures from the first draw in res.

equation_editor

Build production, inefficiency, and risk equations for SFA/MSF

Description

Constructs formula objects for the production function, the inefficiency function, and the production risk function used in stochastic frontier and meta-stochastic frontier estimations. The function takes a generic functional form (FXN) and conditionally augments it with intercept and slope shifters, as well as covariates for the inefficiency (u) and noise (v) equations.

Usage

```
equation_editor(
  data,
  FXN,
  outcome,
  slope_shifter = "NONE",
  intercept_shifters = NULL,
  ulist = NULL,
  vlist = NULL
)
```

Arguments

data	Data.frame or similar object containing all variables referenced in FXN, outcome, slope_shifter, intercept_shifters, ulist, and vlist.
FXN	Functional form of the production function. Typically a named object where FXN[[1]] is a character string representing the production frontier in terms of input variables (e.g., " $\ln I_1 + \ln I_2 + I(1/2 * \ln I_1^2) + \ln I_1 : \ln I_2$ " for a translog).
outcome	Character scalar giving the dependent variable name for the production equation (e.g., "Y" or " $\ln Y$ ").
slope_shifter	Character scalar naming a variable used as a slope shifter in the production function. If equal to "NONE" (default), no slope shifter is applied. Otherwise, the production equation is specified as an interaction between factor(slope_shifter) and FXN[[1]], net of the baseline FXN[[1]].
intercept_shifters	<p>Optional list of intercept shifters for the production function with elements:</p> <ul style="list-style-type: none"> • Svarlist: character vector of continuous shifter variables. • Fvarlist: character vector of categorical shifter variables (included as factor() terms).

If `intercept_shifters` is `NULL` or both components are `NULL`, the production function only uses `FXN[[1]]`.

- `ulist` Optional list of covariates for the inefficiency (u) equation with elements:
- `Svarlist`: character vector of continuous covariates.
 - `Fvarlist`: character vector of categorical covariates (included as `factor()` terms).
- If `ulist` is `NULL` or both components are `NULL`, the inefficiency equation is set to `NULL`.
- `vlist` Optional list of covariates for the production risk (v) equation with elements:
- `Svarlist`: character vector of continuous covariates.
 - `Fvarlist`: character vector of categorical covariates (included as `factor()` terms).
- If `vlist` is `NULL` or both components are `NULL`, the risk equation is set to `NULL`.

Details

Continuous covariates are only included when their coefficient of variation exceeds a small threshold ($|CV| > 0.001$), and categorical covariates are included only if they have more than one level and each level has at least a minimal share of the sample ($> 0.01\%$ of observations). This helps avoid numerical issues from nearly-constant or extremely sparse regressors.

The function builds three formula objects:

Production function (prodfxn) Starts from `outcome ~ FXN[[1]]`. If `intercept shifters` are supplied, they are appended as additive terms (continuous variables directly, categorical variables as `factor()`). When a `slope_shifter` is specified (not "NONE"), the main production part is reparameterized as:

```
outcome ~ factor(slope_shifter) * (FXN[[1]]) - (FXN[[1]]) + shifters
```

Inefficiency function (uequ) Built as $\sim 1 + \dots$ using continuous and/or categorical variables specified in `ulist`. If neither `Svarlist` nor `Fvarlist` are supplied (or all are filtered out), `uequ` is set to `NULL`.

Risk function (vequ) Built analogously to `uequ` using `vlist`. If no valid covariates remain, `vequ` is set to `NULL`.

In all three parts, continuous and categorical candidate variables are screened to avoid near-constant or extremely sparse regressors:

- Continuous variables are retained only if $|sd(x)/mean(x)| > 0.001$.
- Categorical variables are retained only if they have more than one level and the smallest category represents at least 0.01% of the sample.

Value

A list with three components:

- `prodfxn`: a [formula](#) for the production frontier.
- `uequ`: a [formula](#) for the inefficiency equation, or `NULL` if no inefficiency covariates were retained.
- `vequ`: a [formula](#) for the production risk equation, or `NULL` if no risk covariates were retained.

See Also

Other frontier analysis: `fit_organizer()`, `msf_workhorse()`, `sf_functional_forms()`, `sf_model_specifications()`, `sf_workhorse()`, `sfaR_summary()`

`fit_organizer`

Organize frontier coefficients and variance-covariance matrix

Description

Extracts and structures the coefficients and variance-covariance matrix from a fitted stochastic frontier (or related) model, mapping them into a generic parameterization compatible with flexible functional forms such as Cobb-Douglas, translog, quadratic, and linear frontiers.

Usage

```
fit_organizer(fit, number_of_inputs, FXN)
```

Arguments

`fit` Fitted model object (typically the result of a call to `sfacross()` or a similar frontier regression). The object must support `coef()` and `vcov()` methods.

`number_of_inputs` Integer. Number of input variables in the production function. This determines how many first-order and second-order terms are expected.

`FXN` Named object describing the functional form of the production function. The name (i.e., `names(FXN)`) is used to branch between supported forms:

- "TL": translog.
- "QD": quadratic.
- "CD": Cobb-Douglas.
- "LN": linear.

For "TL" and "CD", inputs are expected to enter in logged form as "lnI1", "lnI2", ...; for "QD" and "LN", in levels as "I1", "I2",

Details

For translog and quadratic specifications, the function identifies and orders both first-order and second-order (interaction) terms. For Cobb-Douglas and linear forms, it augments the estimated coefficients with zero-valued second-order terms and builds a conformable variance-covariance matrix (filled with zeros for the non-estimated interactions), so that downstream routines can work with a common parameter vector and matrix.

The function constructs three key objects:

- `est_list`: character vector of coefficient names as they appear in the fitted model object (e.g., "(Intercept)", "lnI1", "lnI2", interaction terms such as "I(1/2 * lnI1 * lnI1)" or "lnI1:lnI2", etc.).
- `est_name`: generic parameter labels corresponding to the structural coefficients, such as "a_0", "a_1", ... for first-order terms and "b_i_j" for second-order terms.
- `est_coef` and `est_vcov`: numeric vector of coefficients and conformable variance-covariance matrix, reordered and expanded (for CD/LN) to match the generic parameterization in `est_name`.

For Cobb-Douglas and linear specifications ("CD", "LN"), only first-order terms are typically estimated directly. The function therefore:

1. Copies the estimated coefficients for first-order terms.
2. Appends zeros for all second-order coefficients so that the length matches the full parameterization implied by `number_of_inputs`.
3. Builds a larger variance-covariance matrix with the original submatrix in the upper-left corner and zeros elsewhere.

This unified representation is useful for elasticities, curvature checks, and minimum-distance procedures that require a full set of parameters and a matching variance-covariance matrix regardless of the underlying functional form.

Value

A list with components:

- `est_coef`: Numeric named vector of organized coefficients. Names follow the generic parameter labels `a_0, a_1, ..., b_i_j`.
- `est_vcov`: Square variance-covariance matrix corresponding to `est_coef`, with row/column names matching the coefficient names.
- `est_list`: Character vector of coefficient names as they were extracted from `fit` (i.e., the original model coefficient labels used to build `est_coef` and `est_vcov`).

See Also

Other frontier analysis: `equation_editor()`, `msf_workhorse()`, `sf_functional_forms()`, `sf_model_specification()`, `sf_workhorse()`, `sfaR_summary()`

`get_crop_area_list` *Extract Matching Crop Area Columns from a Dataset*

Description

Identifies and returns the column names in a dataset corresponding to crop area variables for a specified set of crops. The function looks for column names that start with "Area_" and match any of the crops provided in `selected_crops`.

Usage

```
get_crop_area_list(
  data,
  selected_crops = c("Beans", "Cassava", "Cocoa", "Cocoyam", "Maize", "Millet", "Okra",
    "Palm", "Peanut", "Pepper", "Plantain", "Rice", "Sorghum", "Tomatoe", "Yam")
)
```

Arguments

<code>data</code>	A <code>data.frame</code> or <code>data.table</code> containing crop-related variables. Column names are expected to include fields prefixed with "Area_", such as "Area_Maize".
<code>selected_crops</code>	A character vector specifying the crop names to filter. Defaults to a common set of crops including "Beans", "Cassava", "Cocoa", "Cocoyam", "Maize", "Millet", "Okra", "Palm", "Peanut", "Pepper", "Plantain", "Rice", "Sorghum", "Tomatoe", and "Yam".

Value

A character vector containing the names of columns in data that correspond to the specified crop area variables (e.g., "Area_Maize", "Area_Rice"). Returns an empty vector if no matching columns are found.

get_household_data	<i>Download and Load Harmonized Household Data from the GHAGricProductivityLab Repository</i>
--------------------	---

Description

Retrieves a harmonized household- or farm-level dataset from the **GHAGricProductivityLab** GitHub repository using **piggyback**, stores it in a package-specific cache directory, and returns the dataset as a `data.frame`. The file is downloaded only once and reused from the local cache on future calls.

Usage

```
get_household_data(
  dataset = "harmonized_crop_farmer_data",
  github_token = NULL
)
```

Arguments

dataset	Character string. Base name of the dataset to retrieve (without file extension). Must correspond to a <code>.dta</code> file in the <code>hh_data</code> GitHub release (e.g., "harmonized_crop_farmer_data")
github_token	Optional GitHub personal access token (PAT). If <code>NULL</code> , the function checks the environment variable <code>GHPProdLab_TOKEN</code> . If that is also missing, the <code>piggyback</code> download will use default authentication behavior.

Details

The function downloads a Stata `.dta` file associated with the chosen dataset from the GitHub release labeled `hh_data`. It uses the package-specific cache directory determined by:

```
tools::R_user_dir("GHAGricProductivityLab", which = "cache")
```

If the file already exists locally, it is not re-downloaded.

GitHub Authentication

- If `github_token` is supplied, it is used.
- Otherwise, the function looks for environment variable `GHPProdLab_TOKEN`.
- If neither is available, the function falls back to default GitHub credentials (e.g., from `gh` CLI or cached credentials).

Value

A `data.frame` containing the requested harmonized dataset.

harmonized_data_prep *Prepare Data for Agricultural Productivity Analysis*

Description

Cleans and transforms a dataset by creating new variables, applying log transformations, converting selected variables to factors or characters, and recoding education levels. The function is designed to standardize inputs for further analysis of agricultural productivity.

Usage

```
harmonized_data_prep(data)
```

Arguments

data A `data.frame` or `data.table` containing household- and farm-level variables such as weights, demographic information, and agricultural inputs.

Value

A cleaned and transformed `data.frame` or `data.table` with additional variables ready for analysis.

match_sample_specifications
Draw / Match Sample Specifications

Description

Generates (i) a draw list by sampling EaId within each unique Survey group and (ii) a grid of matching specifications for each bootstrap draw.

Usage

```
match_sample_specifications(number_of_draws = NULL, data, myseed = NULL)
```

Arguments

number_of_draws	Integer. The number of draws to perform per Survey.
data	A <code>data.frame</code> or <code>data.table</code> containing at least the columns Survey and EaId.
myseed	Integer. Seed for random number generation (default 03242025).

Details

Draw list: For each unique value of Survey, the function samples number_of_draws EaId values with replacement and prepends a 0 row (ID = 0) for the baseline. The result is then spread to wide format with one column per Survey.

Matching specs: For each draw ID, creates a set of matching model specifications that include:

- Nearest neighbor with distances: "euclidean", "scaled_euclidean", "mahalanobis", "robust_mahalanobis".
- Nearest neighbor with distance "glm" and links: "logit", "probit", "cloglog", "cauchit".

An ARRAY index is added for convenience.

Value

A list with two elements:

m.specs A data.frame of matching specifications with columns boot, method, distance, link, ARRAY.

drawlist A data.frame in wide format where each Survey is a column and rows correspond to draw ID (0:number_of_draws).

Note

Requires that data contain Survey and EaId. `tidy::spread()` is used for wide reshaping (consider `tidy::pivot_wider()` in new code).

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `treatment_effect_calculation()`, `treatment_effect_summary()`, `write_match_formulas()`

Description

Performs Meta Stochastic Frontier (MSF) analysis in several stages:

1. **Naive SF:** Fits a pooled stochastic frontier and computes baseline technical efficiency (TE) measures.
2. **Group SF:** If a technology variable tvar is supplied, splits the sample into technology groups and fits group-specific SF models.
3. **Meta SF (unmatched and matched):** Constructs a meta-frontier using fitted values from group models and, optionally, nearest-neighbor matching specifications. From this, it derives technology gap ratios (TGR) and meta-technical efficiency (MTE).
4. **Summaries and distributions:** Produces weighted and unweighted summaries and distributional statistics for efficiency, elasticity, and risk across technologies, samples (unmatched/matched), and surveys.

Usage

```
msf_workhorse(
  data,
  yvar,
  xlist,
  ulist = NULL,
  vlist = NULL,
  wvar = NULL,
  slope_shifter = "NONE",
  intercept_shifters = NULL,
  f,
  d,
  identifiers,
  include_trend = FALSE,
  tvar = NULL,
  matching_type = NULL,
  ulistM = NULL,
  intercept_shiftersM = NULL
)
```

Arguments

data	A data.frame or data.table containing the estimation sample. Must include the dependent variable yvar , the inputs in xlist , the weight variable wvar , unique ID variables in identifiers , any inefficiency and risk covariates in ulist and vlist , the technology variable tvar (if used), and any variables required by matching objects on disk (e.g., "Surveyx", "EaId", "HhId", "Mid", "unique_identifier").
yvar	Character scalar. Name of the dependent (output) variable used in the production frontier.
xlist	Character vector of input variable names (e.g., land, labor, capital) passed through to sf_workhorse .
ulist	Optional named list specifying variables in the inefficiency function for the naive and group frontiers. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> .
vlist	Optional named list specifying variables in the production risk (noise) function for the naive and group frontiers. Same structure as ulist . If <code>NULL</code> , a homoskedastic noise term is usually implied.
wvar	Character scalar. Name of the sampling/observation weight variable in data . Observations with zero weight are dropped.
slope_shifter	Character scalar giving the name of a slope-shifter variable in data (e.g., technology shifter, policy dummy). Defaults to "NONE" for no slope shifter.
intercept_shifters	Optional named list of intercept shifter variables for the naive and group frontiers. Typical structure: <code>list(Svarlist = c(...), Fvarlist = c(...))</code> , passed to sf_workhorse .
f	Functional form index for the production frontier (e.g., Cobb-Douglas, translog, quadratic). The index is passed to sf_workhorse and ultimately to the internal form-selection utilities (e.g., <code>sf_functional_forms</code>).
d	Distributional form index for the inefficiency term (e.g., half-normal, truncated-normal). Passed to sf_workhorse .

<code>identifiers</code>	Character vector of variable names that uniquely identify observations (e.g., c("unique_identifier", "Survey", "CropID", "HhId", "EaId", "Mid")). These IDs are used when merging scores and summarizing by technology or sample.
<code>include_trend</code>	Logical; if TRUE, the last element in <code>xlist</code> is treated as a trend/technology variable in the production function. Passed through to <code>sf_workhorse</code> . Defaults to FALSE.
<code>tvar</code>	Optional character scalar naming the technology variable (e.g., region, system, period) used to define groups for group SF and meta-frontier estimation. If NULL, no technology groups are formed and only the naive frontier and TE are computed.
<code>matching_type</code>	Optional character scalar controlling which nearest-neighbor matching specifications to use for meta-frontier estimation. When not NULL, the function reads matching specifications and matched samples from "results/matching/" and related RDS files. Setting <code>matching_type</code> = "optimal" restricts attention to the subset of matching specifications labeled as optimal.
<code>ulistM</code>	Optional named list specifying inefficiency-function covariates for the meta-frontier (matched/unmatched) estimation. If NULL, defaults to <code>ulist</code> .
<code>intercept_shiftersM</code>	Optional named list of intercept shifters for the meta-frontier estimation. If NULL, defaults to <code>intercept_shifters</code> .

Details

The function is designed as a high-level orchestrator around `sf_workhorse`, assembling naive, group-level, and meta-frontier results into a unified output structure suitable for MSF analysis.

The workflow can be summarized as:

- **Naive SFA (TE0):** Calls `sf_workhorse` on the full sample to obtain baseline efficiencies (teBC, teJLMS, teMO). If no technology variable is specified (`tvar` = NULL), these naive efficiencies are the final scores and only TE-related summaries are produced.
- **Group SFA (TE):** When `tvar` is provided, the sample is recoded into numeric technology groups (Tech), and `sf_workhorse` is applied separately to each group. Group-level efficiencies are combined into TE measures by group and survey.
- **Meta-frontier (TGR and MTE):** Using fitted group-frontier values (\hat{Y}) from the restricted models, the function fits meta-frontiers (unmatched and, optionally, matched) again via `sf_workhorse`. Technology gap ratios (TGR) are derived from these meta-frontiers, and meta-technical efficiency (MTE) is computed as $TE * TGR$.
- **Matching layer (optional):** If `matching_type` is given, the function loads matching specifications (`mspecs`, `mspecs_optimal`) and corresponding matched samples from disk (e.g., "results/matching/"), fits additional meta-frontiers by matching specification, and augments the set of scores and summaries with these matched samples.
- **Summaries and distributions:**

Scores: TE0, TE, TGR, and MTE are reshaped into long form and aggregated to compute weighted/unweighted means, medians, modes, and distribution histograms (both count-based and weight-based) across surveys, samples, technologies, and estimation types. When `tvar` is not NULL, the function also computes technology gaps (level and percent) relative to the minimum technology in `technology_legend`.

Elasticities: Elasticities from the underlying SF models (naive, group, meta) are combined and summarized in a similar way, including technology-gap metrics for elasticities when `tvar` is provided.

Risk: Risk measures derived from `sf_workhorse` are combined, summarized, and used to build distributional statistics analogous to those for efficiency.

- **LR tests:** When `tvar` is not `NULL`, the function builds likelihood ratio test statistics comparing a naive pooled frontier to the combination of group and meta-frontiers, storing these as rows with `CoefName = "LRT"` in the `sf_estm` output.

Throughout, the function uses **dplyr**, **tidyr**, **data.table**, **doBy**, and related helper functions for reshaping and summarizing the output of `sf_workhorse`.

Value

A named list with the following components:

- `sf_estm`: Data.frame of coefficient-level summaries from all naive, group, and meta-frontiers (including LR tests), augmented with technology identifiers (Tech) and sample labels (e.g., "unmatched", matching-spec names).
- `ef_mean`: Data.frame of aggregated efficiency statistics (TE0, TE, TGR, MTE), including weighted/unweighted means, medians, and modes, by survey, sample, technology, type, estimation type, and restriction status. When `tvar` is supplied, also includes efficiency gap levels and percentages.
- `ef_dist`: Data.frame of efficiency distributions (histogram counts and weights) over efficiency ranges, by survey, sample, technology, type, estimation type, and restriction.
- `rk_dist`: Data.frame of risk distributions (histogram counts and weights) analogous to `ef_dist`, with `type = "risk"`.
- `el_mean`: Data.frame of aggregated elasticity statistics (weighted/unweighted means, medians, modes) by input, survey, sample, technology, and restriction, including elasticity-gap measures when `tvar` is provided.
- `rk_mean`: Data.frame of aggregated risk statistics (weighted and unweighted) by survey, sample, technology, and restriction, with risk-gap measures when `tvar` is provided.
- `el_samp`: Data.frame of observation-level elasticities for all models and samples (excluding the synthetic "GLSS0" survey rows).
- `ef_samp`: Data.frame of observation-level efficiency scores (TE0, TE, TGR, MTE) including IDs, weights, technologies, and sample labels.
- `rk_samp`: Data.frame of observation-level risk measures for all models and samples (excluding the synthetic "GLSS0" survey rows).

See Also

Other frontier analysis: `equation_editor()`, `fit_organizer()`, `sf_functional_forms()`, `sf_model_specification()`, `sf_workhorse()`, `sfaR_summary()`

Description

Builds a tidy coefficient-and-test table from a fitted stochastic frontier model (typically from **sfaR**). The output combines:

- Maximum likelihood estimates of the frontier and auxiliary parameters (coefficients, standard errors, z-values, p-values), and
- Model-level diagnostics and tests, including variance components, skewness tests for OLS residuals, and a likelihood-ratio test for the presence of inefficiency.

Usage

```
sfaR_summary(fit)
```

Arguments

fit	A fitted stochastic frontier model object from sfaR . The object must support <code>summary()</code> and contain elements such as <code>mlLoglik</code> , <code>olsSkew</code> , <code>olsM3Okay</code> , <code>CoelliM3Test</code> , and <code>AgostinoTest</code> (as produced by sfaR fits).
------------	---

Details

This function is intended as a convenient post-estimation summarizer whose output can be merged with other model summaries or used directly in tables and figures.

The function proceeds in several steps:

1. Extracts the maximum-likelihood coefficient table from `summary(fit)$mlRes` and renames columns to `Estimate`, `StdError`, `Zvalue`, `Pvalue`.
2. Constructs a one-row data.frame (TEST) containing model-level statistics returned by `summary(fit)`, such as: `Nobs`, `nXvar`, `nuZUvar`, `nvZVvar`, `mlLoglik`, `AIC`, `BIC`, `HQIC`, `sigmavSq`, `sigmauSq`, `Varu`, `Eu`, and `Expu`. From these, it computes:
 - `Sigma = sqrt(sigmauSq + sigmavSq)`, and
 - `Gamma = sigmauSq / (Sigma^2)`.

It also overwrites `mlLoglik` with `fit$mlLoglik`, adds OLS skewness (`olsSkew`), a flag for the expected skewness (`olsM3Okay`), and reshapes everything into a `CoefName/Estimate` table with `Pvalue = NA`.

3. Appends additional tests:

Coelli's M3 test Stored as `CoefName = "CoelliM3Test"` with statistic and p-value taken from `fit$CoelliM3Test`.

D'Agostino tests Three rows are added for "`AgostinoOmn`", "`AgostinoSkw`", and "`AgostinoKrt`", using `fit$AgostinoTest@test$statistic` and `fit$AgostinoTest@test$p.value`.

LR test of inefficiency One row with `CoefName = "LRInef"`, where the test statistic is `summary(fit)$chisq`. The p-value is computed using a chi-bar-square distribution via `emdbook::qchibarsq()` at 0.90, 0.95, and 0.99 quantiles, and then discretized to 0.10, 0.05, 0.01, or 1.

Finally, the ML coefficient table and the TEST rows are stacked (using `data.table::rbindlist()`) into a single data.frame with consistent columns.

Value

A data.frame with at least the following columns:

- CoefName: Name of the parameter or diagnostic/statistic.
- Estimate: Point estimate or test statistic.
- StdError: Standard error (for ML coefficients; NA for many diagnostics).
- Zvalue: z-statistic for ML coefficients (NA otherwise).
- Pvalue: p-value for ML coefficients and tests when available (NA if not applicable).

See Also

Other frontier analysis: [equation_editor\(\)](#), [fit_organizer\(\)](#), [msf_workhorse\(\)](#), [sf_functional_forms\(\)](#), [sf_model_specifications\(\)](#), [sf_workhorse\(\)](#)

sf_functional_forms *Generate Functional and Distribution Forms for MSF Models*

Description

Builds algebraic strings for common production-function specifications and assembles a catalog of distributional assumptions for one-sided inefficiency terms used in stochastic frontier / multi-stage frontier (MSF) analysis.

Usage

```
sf_functional_forms(number_of_inputs = 5, include_trend = FALSE)
```

Arguments

number_of_inputs	Integer. Number of input variables to enumerate in the functional forms (creates symbols $I_1 \dots I_k$ and $\ln I_1 \dots \ln I_k$). Default 5.
include_trend	Logical. If TRUE, adjusts the Transcendental production form (TP) by dropping the last linear input term. See Note regarding when TP is present. Default FALSE.

Details

Functional forms (returned as character strings):

- CD: Cobb-Douglas - sum of logs of inputs, e.g., $\ln I_1 + \ln I_2 + \dots + \ln I_k$.
- TL: Translog - CD plus all second-order log terms (squares and cross-products), e.g. %
- LN: Linear (commented out in current code). %
- QD: Quadratic (commented out). %
- GP: Generalized (commented out). %
- TP: Transcendental (commented out; see Note).

Terms are generated programmatically for $k = \text{number_of_inputs}$, with: I_i denoting input levels and $\ln I_i$ their logarithms.

Distributional forms (list of lists): Each entry is named and contains `list(name, scaling = <logical>)`. Included options: `hnnormal`, `tnormal`, `exponential`, `tslaplace`, `genexponential`, `tnormal_scaled`, `rayleigh`, `uniform`, `gamma`, `lognormal`, `weibull`. The scaling flag indicates whether the Wang-Schmidt (2002) scaling property is used in that specification.

Value

A list with two components:

- `fxnforms`: named list of character strings for functional forms (e.g., `$CD`, `$TL`). Forms that are commented out in the implementation will not appear.
- `distforms`: named list of distribution specifications; each is a list with elements `[1]` = distribution name and `scaling` = logical.

Note

The current implementation defines `CD` and `TL` and removes `NULL` entries before returning. Other forms (`LN`, `QD`, `GP`, `TP`) are commented out. If you plan to set `include_trend = TRUE`, ensure `TP` is actually included in `fxnforms`; otherwise, modifying `fxnforms$TP` will not have any effect (and would error if `TP` is not defined).

See Also

Other frontier analysis: `equation_editor()`, `fit_organizer()`, `msf_workhorse()`, `sf_model_specifications()`, `sf_workhorse()`, `sfaR_summary()`

sf_model_specifications

Build Model Specifications

Description

Construct a specification table for frontier analysis by combining alternative production-function forms and distributional assumptions with technology choices and disaggregation levels. The function first creates a pooled grid of functional forms (`fxnforms`) and distributions (`distforms`), then augments it with crop-specific and demographic splits, and finally expands technology choices.

Usage

```
sf_model_specifications(
  data,
  distforms = sf_functional_forms()$distforms,
  f xnforms = sf_functional_forms()$fxnforms,
  technology_variables,
  mainF = 2,
  mainD = 1,
  demographic_variables = c("Female", "Region", "Ecozon", "EduCat", "EduLevel", "AgeCat"),
  crop_list = c("Beans", "Cassava", "Cocoa", "Cocoyam", "Maize", "Millet", "Okra",
    "Palm", "Peanut", "Pepper", "Plantain", "Rice", "Sorghum", "Tomatoe", "Yam")
)
```

Arguments

<code>data</code>	data.frame or data.table. A dataset that (ideally) contains the columns named in <code>demographic_variables</code> . Note: The current implementation references a global object <code>DATA</code> (not <code>data</code>) when deriving demographic levels; ensure <code>DATA</code> exists with those columns or adapt your environment accordingly.
-------------------	---

distforms named list of distributional forms. If NULL, defaults to `functional_forms()$distforms`.
fxnforms named list of functional forms. If NULL, defaults to `functional_forms()$fxnforms`.
technology_variables character vector of technology variables. The first element is taken as the default/primary technology.
mainF integer index of the preferred functional form (in `fxnforms`). Default is 2.
mainD integer index of the preferred distribution (in `distforms`). Default is 1.
demographic_variables character vector of column names used for demographic disaggregation. Defaults to `c("Female", "Region", "Ecozon", "EduCat", "EduLevel", "AgeCat")`.
crop_list character vector of crop names for crop-level disaggregation. Defaults to a selection of common crops.

Details

Procedure (high level):

1. Build a pooled grid of all `fxnforms * distforms`.
2. Keep rows matching `mainF` and `mainD` and tag pooled specs.
3. Add crop-specific rows for `crop_list` (using the main form/dist).
4. Add demographic splits for each variable in `demographic_variables`, using unique levels in `DATA[, var]` (see note in *Warning*).
5. Set `TechVar` to the first element of `technology_variables`; if additional technology variables are provided, add pooled rows for them with the main form/dist.
6. Reorder to prioritize `mainD` and create an `matching_type` flag with values "fullset" and "optimal".
7. Remove demographic rows for the default demographic list and retain only `level == "Pooled"` in the final output.

Returned columns (final table):

- `disasg` - disaggregation variable (e.g., "CropID").
- `level` - disaggregation level (e.g., "Pooled" or a crop/level).
- `TechVar` - selected technology variable.
- `f` - index of functional form (position in `fxnforms`).
- `d` - index of distribution (position in `distforms`).
- `matching_type` - marker column ("fullset" / "optimal").

Value

A `data.frame` (data.table-compatible) of model specifications with columns `disasg`, `level`, `TechVar`, `f`, `d`, and `matching_type`.

Warning

The function calls `unique(DATA[, w])` inside a `tryCatch`; this requires a global object `DATA` containing the demographic columns. If `DATA` is absent or lacks a given variable, that split is skipped.

Assumptions

- `technology_variables` is non-empty; its first element is the default.
- `mainF` and `mainD` are valid indices within `fxnforms` and `distforms`, respectively.

See Also

Other frontier analysis: `equation_editor()`, `fit_organizer()`, `msf_workhorse()`, `sf_functional_forms()`, `sf_workhorse()`, `sfaR_summary()`

`sf_workhorse`

Stochastic frontier workhorse (unrestricted and shape-constrained)

Description

Fits a stochastic frontier model for a given production specification and, when monotonicity is sufficiently violated, re-estimates a shape-constrained (restricted) frontier using minimum-distance methods. The function wraps the underlying sfaR estimation routines and a set helper utilities (e.g., `sf_functional_forms()`, `equation_editor()`, `Fxn.fit_organizer()`, `translogEla()`, curvature/monotonicity checks) to produce:

1. Unrestricted stochastic frontier estimates and diagnostics.
2. Shape-constrained estimates that enforce monotonicity (and related regularity conditions) when needed.
3. Observation-level efficiencies, fitted values, and risk measures.
4. Observation-level elasticities, including a "returns-to-scale" term.

Usage

```
sf_workhorse(
  data,
  yvar,
  xlist,
  wvar = NULL,
  d,
  f,
  slope_shifter = "NONE",
  intercept_shifters = NULL,
  ulist = NULL,
  vlist = NULL,
  identifiers,
  include_trend = FALSE
)
```

Arguments

<code>data</code>	A <code>data.frame</code> or <code>data.table</code> containing all variables required for estimation, including the dependent variable <code>yvar</code> , inputs <code>xlist</code> , the weight variable <code>wvar</code> (if used), the unique ID variables listed in <code>identifiers</code> , and any variables referenced in <code>slope_shifter</code> , <code>intercept_shifters</code> , <code>ulist</code> , or <code>vlist</code> .
-------------------	---

yvar	Character scalar. Name of the dependent/output variable in data used for the production frontier.
xlist	Character vector of input variable names (e.g., land, labor, capital). The order of variables in xlist determines how they map into the generic input labels I1, I2, ... used in the functional-form utilities.
wvar	Optional character scalar. Name of the sampling/observation weight variable. If NULL, all observations are given unit weight.
d	Distributional form index for the inefficiency term. This is passed to sf_functional_forms() and ultimately to sfacross() (e.g., to select half-normal, truncated-normal, etc.). The exact mapping is determined by sf_functional_forms().
f	Functional form index for the production frontier (e.g., Cobb-Douglas, translog, quadratic), used to select from the list returned by sf_functional_forms(). The name of the chosen form (e.g., "CD", "TL", "QD", "LN") influences logging of variables and regularity checks.
slope_shifter	Character scalar giving the name of a slope-shifter variable in data (e.g., technology shift, policy dummy). Defaults to "NONE" for no slope shifter and is passed to equation_editor().
intercept_shifters	Optional named list of intercept shifter variables for the production function. Typical structure: list(Svarlist = c(...), Fvarlist = c(...)), where the specific interpretation is handled inside equation_editor().
ulist	Optional named list describing the inefficiency-function covariates. Typical structure: list(Svarlist = c(...), Fvarlist = c(...)), passed as uhet and muhet to sfacross() when non-NULL.
vlist	Optional named list describing the production-risk (noise) covariates. When non-NULL, used as vhet in sfacross() to allow heteroskedasticity of the noise term.
identifiers	Character vector of variable names that uniquely identify observations (e.g., c("unique_identifier", "Survey", "CropID", "HhId", "EaId", "Mid")). These IDs are carried through to efficiency, elasticity, and risk outputs.
include_trend	Logical; if TRUE, the last element in xlist is treated as a trend/technology variable and handled accordingly in the functional-form and elasticity calculations. Defaults to FALSE.

Details

The function proceeds in several steps:

- 1. Setup and functional form selection:** Using sf_functional_forms(), the function determines the appropriate production functional form and distributional specification based on the number of inputs (number_of_inputs), the f and d indices, and the include_trend flag. This includes whether the model is specified in logs (e.g., CD/TL/GP/TP) or levels.
- 2. Variable construction and equation building:** Inputs in xlist are mapped to generic labels (I1, I2, ...) and their log transforms are created when needed. The outcome variable is set to Y or lnY. The function then calls equation_editor() to construct the production, inefficiency, and risk equations used by sfacross().
- 3. Unrestricted stochastic frontier estimation:** The function iterates over a grid of optimization methods (nr, nm, bfgs, bhhh, ...) and tolerance settings, calling sfacross() until a successful convergence is reported. It extracts observation-level efficiencies via sfaR::efficiencies(), constructs fitted values, and builds an observation-level risk metric based on squared deviations from mean output.

4. **Elasticities and regularity checks:** Using `Fxn.fit_organizer()` and `translogEla()`, the function computes elasticities and returns-to-scale-type measures, and evaluates monotonicity and curvature via `translogCheckMono()` and `translogCheckCurvature()` (or a coefficient-sign check for CD/LN forms).
5. **Shape-constrained frontier (if needed):** If the monotonicity measure `mono` falls below 0.80, the function constructs a set of linear restrictions using `translogMonoRestr()` and solves a quadratic programming problem (via `quadprog::solve.QP` and fall-back matrix inversions using **Matrix**, **MASS**, **corpcor**) to obtain constrained coefficients. For TL/QD forms a constrained frontier (`lcFitted`) is computed with `translogCalc()`; for CD/LN forms, constrained coefficients are obtained via a SEM representation using **lavaan**.
6. **Re-estimation under constraints:** Given the constrained frontier, the function rebuilds the production equation with `equation_editor()` and re-estimates a stochastic frontier (`sfc`) using the same optimization grid. Constrained efficiencies, risks, and elasticities are computed and summarized, and regularity checks are repeated.

Unrestricted and restricted summaries are then stacked, with a `restrict` flag ("Unrestricted" vs. "Restricted"), for downstream comparison.

Value

A named list with four data.frames:

- `sf`: Coefficient-level results combining unrestricted and, when applicable, restricted SFA estimates, plus monotonicity/curvature diagnostics. Columns include (at minimum) `CoefName`, `Estimate`, `StdError`, `Zvalue`, `Pvalue`, and `restrict`.
- `ef`: Observation-level efficiency results (unrestricted and restricted), including the ID variables in `identifiers`, weights, efficiency measures (e.g., `u`), model-fitted values (`mlFitted`), and `restrict`.
- `el`: Observation-level elasticities and returns-to-scale-type measures. Includes IDs, weights, elasticity columns (e.g., `el1`, `el2`, ..., and the summed elasticity), and `restrict`.
- `rk`: Observation-level risk measures (unrestricted and restricted), including IDs, weights, and risk, plus `restrict`.

See Also

Other frontier analysis: `equation_editor()`, `fit_organizer()`, `msf_workhorse()`, `sf_functional_forms()`, `sf_model_specifications()`, `sfaR_summary()`

Description

Initialize Study Environment and Directory Structure

Usage

```
study_setup(
  myseed = 1980632,
  project_name,
  local_directories = list(home = file.path("replications", project_name), output =
    file.path("replications", project_name, "output"), matching =
    file.path("replications", project_name, "output", "matching"), treatment_effects =
    file.path("replications", project_name, "output", "treatment_effects"), estimations =
    file.path("replications", project_name, "output", "estimations"), figure_data =
    file.path("replications", project_name, "output", "figure_data"), figure =
    file.path("replications", project_name, "output", "figure")))
)
```

Arguments

- myseed** Numeric/integer scalar seed. Will be coerced to integer. Default: 1980632.
project_name Length-1, non-NA character project name (required).
local_directories Named list of character paths to create. Defaults use `project_name`.

Value

List with `wd` (directories) and `seed`.

treatment_effect_calculation

Compute log-linear treatment effects (ATE/ATET/ATEU) for multiple outcomes

Description

For a given matching specification index `i`, this function:

1. Loads the corresponding matched weights file from `matching_output_directory`,
2. Merges those weights into the analysis data via `unique_identifier`,
3. Optionally normalizes each outcome by the first element of `outcome_variables` (e.g., to construct per-hectare or per-unit measures),
4. Fits a weighted log-linear model for each outcome with interactions between `Treat` and the supplied matching covariate formulas,
5. Transforms fitted differences into percentage treatment effects, trims extreme values, and computes weighted ATE, ATET, and ATEU alongside model fit statistics.

Usage

```
treatment_effect_calculation(
  data,
  outcome_variables,
  normalize = NULL,
  i,
```

```

    match_specifications,
    matching_output_directory,
    match_formulas
)

```

Arguments

<code>data</code>	A <code>data.frame</code> containing at least: <ul style="list-style-type: none"> • <code>unique_identifier</code>: unit identifier used for merging matched weights, • <code>Treat</code>: treatment indicator (logical or 0/1), • <code>Area</code>: plot size (positive; used for input checks and typical scaling), • all variables named in <code>outcome_variables</code>, • any covariates referenced in <code>match_formulas\$general_match</code> and <code>match_formulas\$exact_match</code>.
<code>outcome_variables</code>	Character vector of outcome variable names (e.g., <code>c("HrvstKg", "Area", "SeedKg", "HHLaborAE", ...)</code>)
<code>normalize</code>	Logical scalar. If TRUE, each outcome in <code>outcome_variables</code> is divided by the first element of <code>outcome_variables</code> in the merged data (e.g., to obtain outcomes per hectare or per unit).
<code>i</code>	Integer index selecting the row of <code>match_specifications</code> and the corresponding matching result file (zero-padded via <code>ARRAY</code> to "NNNN.rds").
<code>match_specifications</code>	A <code>data.frame</code> with at least a column <code>ARRAY</code> . The <code>i</code> -th row is used to locate the matching output file and is <code>cbind</code> -ed to the returned results.
<code>matching_output_directory</code>	Directory containing per-specification matching results named "match_0001.rds", "match_0002.rds", etc. Each RDS is expected to contain an element <code>\$md</code> with columns <code>unique_identifier</code> and <code>weights</code> .
<code>match_formulas</code>	A list with: <ul style="list-style-type: none"> • <code>general_match</code>: a formula; the RHS is accessed via <code>as.character(...)[3]</code>, • <code>exact_match</code>: a formula; the RHS is accessed via <code>as.character(...)[2]</code>.

Details

For each outcome `oc` in `outcome_variables`, the function:

1. Filters to observations with non-missing `Treat`, the outcome, and weights,
2. Constructs a log-linear model of the form

$$\log(oc + \text{eps}) \sim Treat * (\text{general_match_rhs} + \text{exact_match_rhs})$$

where `eps = min(oc[oc > 0]) * 1/100` in the estimation sample,

3. Fits the model using `lm()` with weights given by the matched weights,
4. Predicts fitted log outcomes for treated (`Treat = 1`) and untreated (`Treat = 0`) counterfactuals,
5. Computes individual percentage treatment effects

$$TE_{OLS} = \{\exp(\hat{y}_1 - \hat{y}_0) - 1\} \times 100,$$

trims them to the interval $[-100, 100]$, and then forms:

- ATE: weighted mean of TE_{OLS} over all units,
- ATET: weighted mean of TE_{OLS} among treated units,
- ATEU: weighted mean of TE_{OLS} among untreated units.

In addition, for each outcome the function records model diagnostics: AIC, log-likelihood, R^2 , adjusted R^2 , F-statistic, and the sample size used in the regression.

Value

A `data.frame` (`atet_scalar`) with the columns from `match_specifications[i,]` repeated across rows, and the additional columns:

- `outcome`: outcome variable name,
- `treatment`: name of the treatment indicator (always "Treat"),
- `level`: one of "ATE", "ATET", "ATEU", "aic", "ll", "R2", "N", "Ft", "R2a",
- `est`: numeric estimate corresponding to level.

One block of rows is returned per outcome in `outcome_variables`. If a model fails for a given outcome, that outcome is skipped.

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `match_sample_specifications()`, `treatment_effect_summary()`, `write_match_formulas()`

`treatment_effect_summary`

Summarize Treatment Effect Estimates Across Matching Specifications

Description

Aggregates and summarizes treatment effect results stored as `.rds` files in a specified output directory. The function reads each RDS file, combines all results into a single data frame, filters out invalid estimates, and computes jackknife-style summary statistics (mean, standard error, sample size, t-value, and p-value) by matching specification and outcome.

Usage

```
treatment_effect_summary(treatment_effects_output_directory)
```

Arguments

`treatment_effects_output_directory`

Character string giving the path to the directory containing `.rds` files with treatment effect results (e.g., outputs of `treatment_effect_calculation()` saved across specifications).

Details

The function proceeds through the following steps:

1. Lists and reads all `.rds` files in `treatment_effects_output_directory`.
2. Combines them into a single long-format data frame using `data.table::rbindlist()`.
3. Removes rows with invalid estimates (NA, NaN, Inf, or -Inf).
4. Separates base (non-bootstrap) results where `boot == 0` and joins these with jackknife summary statistics computed via `doBy::summaryBy()` across method, distance, link, outcome, and level.
5. Renames summary columns to:

- jack_mean: mean estimate
- jack_se: standard error
- jack_n: sample size (number of replicates)
- jack_tvl: t-value (jack_mean / jack_se)
- jack_pvl: two-sided p-value

Value

A `data.frame` containing summarized treatment effect estimates with columns:

- method, distance, link, outcome, level
- est (original estimate for non-bootstrap sample)
- jack_mean, jack_se, jack_n, jack_tvl, jack_pvl

See Also

Other matching: `covariate_balance()`, `draw_matched_samples()`, `match_sample_specifications()`, `treatment_effect_calculation()`, `write_match_formulas()`

`write_match_formulas` *Construct Matching Formulas for Exact and General Matching*

Description

Builds two types of matching formulas commonly used in propensity score or covariate matching workflows:

- An **exact match** formula specifying categorical variables to match exactly.
- A **general match** formula specifying numeric and factor covariates for distance-based matching.

The function returns both formulas as character strings that can be used in matching functions such as `MatchIt::matchit()` or `Matching::Match()`.

Usage

```
write_match_formulas(
  match_variables_exact,
  match_variables_scaler,
  match_variables_factor
)
```

Arguments

<code>match_variables_exact</code>	A character vector of variable names to be included in the exact match component (e.g., "gender", "region").
<code>match_variables_scaler</code>	A character vector of continuous or scalar variable names to include as covariates in the general match formula (e.g., "age", "income").
<code>match_variables_factor</code>	A character vector of factor variable names to be included as dummy-coded categorical covariates in the general match formula.

Value

A named list with two elements:

`exact_match` A string representing the exact match formula (factor variables only).

`general_match` A string representing the general match formula, including both continuous and factor covariates on the right-hand side of `Treat ~`.

See Also

Other matching: [covariate_balance\(\)](#), [draw_matched_samples\(\)](#), [match_sample_specifications\(\)](#), [treatment_effect_calculation\(\)](#), [treatment_effect_summary\(\)](#)

Index

- * **frontier analysis**
 - equation_editor, 9
 - fit_organizer, 11
 - msf_workhorse, 15
 - sf_functional_forms, 20
 - sf_model_specifications, 21
 - sf_workhorse, 23
 - sfaR_summary, 18
- * **matching**
 - covariate_balance, 2
 - draw_matched_samples, 3
 - match_sample_specifications, 14
 - treatment_effect_calculation, 26
 - treatment_effect_summary, 28
 - write_match_formulas, 29
- covariate_balance, 2, 4, 15, 28–30
- draw_matched_samples, 3, 3, 15, 28–30
- draw_msf_estimations, 4
- draw_msf_summary, 7
- equation_editor, 9, 12, 18, 20, 21, 23, 25
- fit_organizer, 10, 11, 18, 20, 21, 23, 25
- formula, 10
- get_crop_area_list, 12
- get_household_data, 13
- harmonized_data_prep, 14
- match_sample_specifications, 3, 4, 14,
28–30
- msf_workhorse, 10, 12, 15, 20, 21, 23, 25
- sf_functional_forms, 10, 12, 18, 20, 20, 23,
25
- sf_model_specifications, 10, 12, 18, 20,
21, 21, 25
- sf_workhorse, 10, 12, 16–18, 20, 21, 23, 23
- sfaR_summary, 10, 12, 18, 18, 21, 23, 25
- study_setup, 25
- treatment_effect_calculation, 3, 4, 15,
26, 29, 30
- treatment_effect_summary, 3, 4, 15, 28, 28,
30
- write_match_formulas, 3, 4, 15, 28, 29, 29